

Midas

On-the-fly-Schema-Migration Tool

For



Schema Migration Problems

- Applications have to hand-roll their own schema migration infrastructure or use some third-party tool
- Difficult to migrate TBs of data without downtime
 - unacceptable from SLA stand-point!
- How about: on-the-fly schema migration - *a Midas Touch?*

Zero-downtime Deployment

- ✦ Expansion Scripts
 - ✦ Apply changes to the documents safely that do not break backwards compatibility with existing version of the application.
 - ✦ e.g Adding, copying, merging, splitting fields in a document.
- ✦ Contraction Scripts
 - ✦ Clean up any database schema that is not needed after the upgrade.
 - ✦ e.g. removing fields from a document.

The Mechanics


- ✦ Run Expansion scripts before upgrading the application
- ✦ Upgrade the cluster, a node at a time
- ✦ Run Contraction Scripts
 - ✦ Once the system has been completely upgraded and deemed stable.
 - ✦ Typically, contractions can be run, say days/weeks after complete validation.

Do we need DB rollback?

- ✦ Short Answer
 - ✦ No
- ✦ Long Answer
 - ✦ Reversing DB changes can lead to potential loss of data or leave it in an inconsistent state.
 - ✦ Its safer to rollback application without needing to rollback DB changes as expansions are backward compatible.


How it works?

An Architectural Overview



How it works?

An Architectural Overview

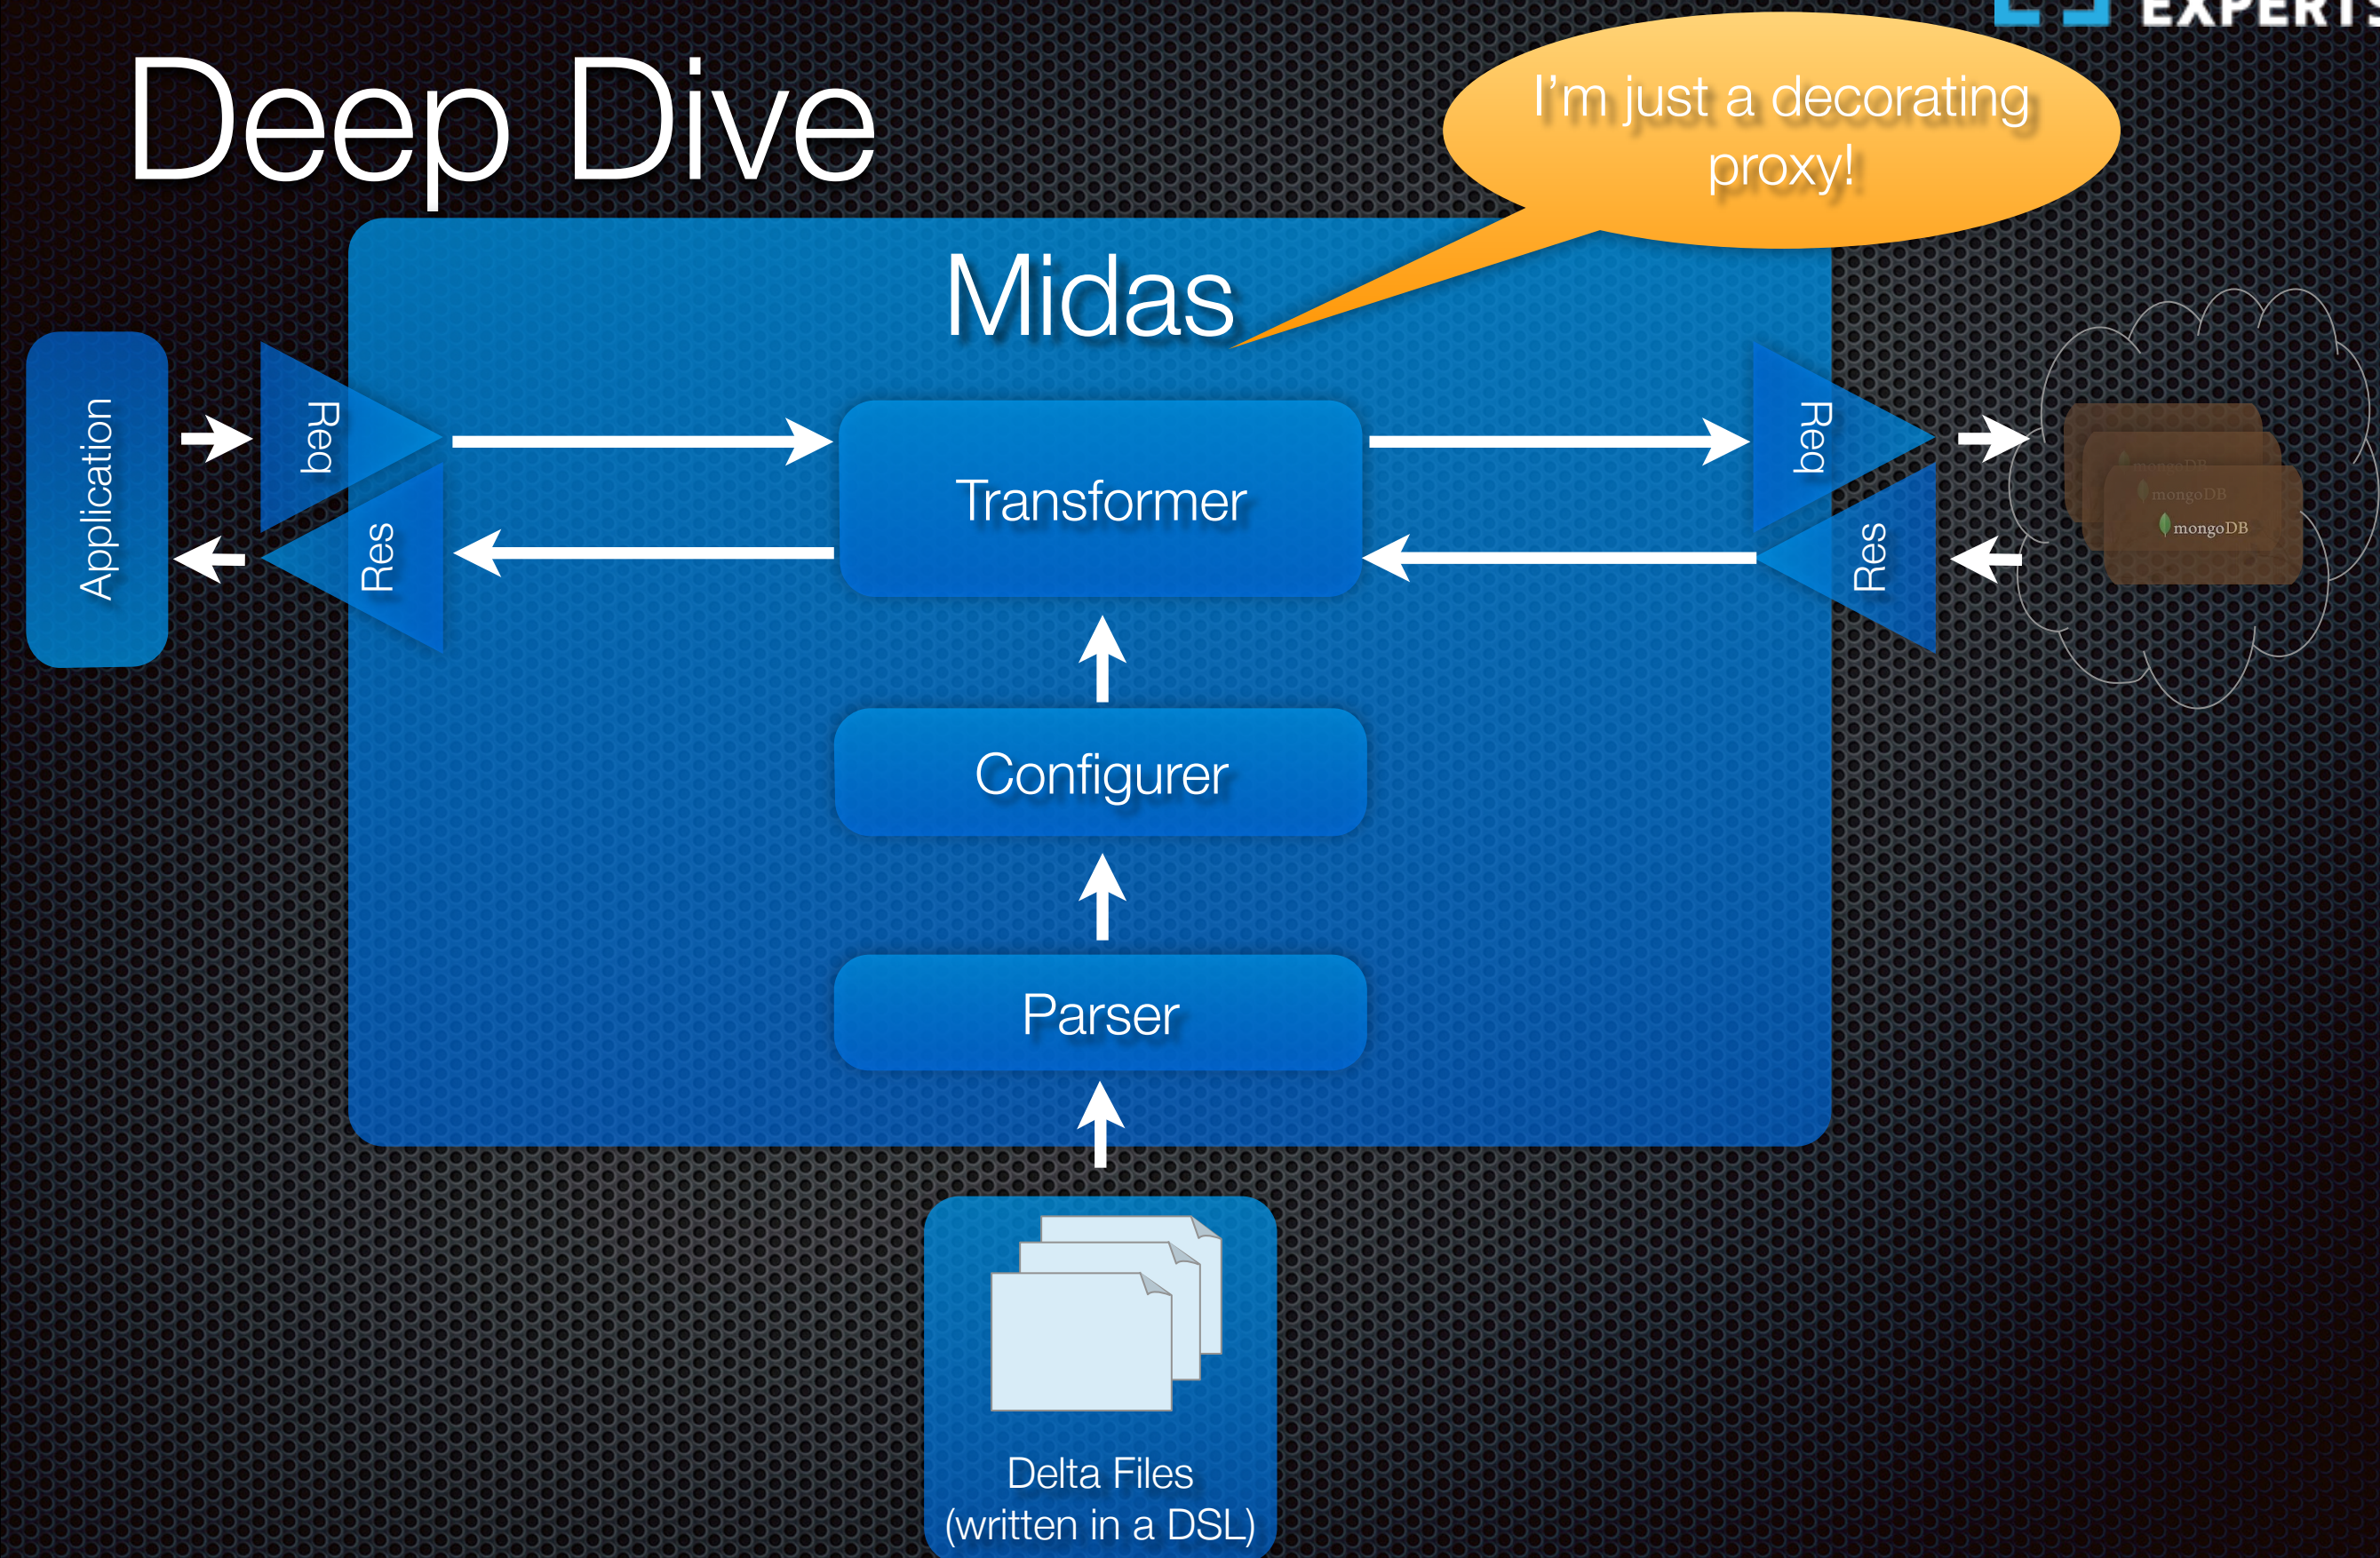


- Intercepts Responses at Protocol Level
- Upgrades/Downgrades Document schema in-transit

Protocol Level brings Transparency

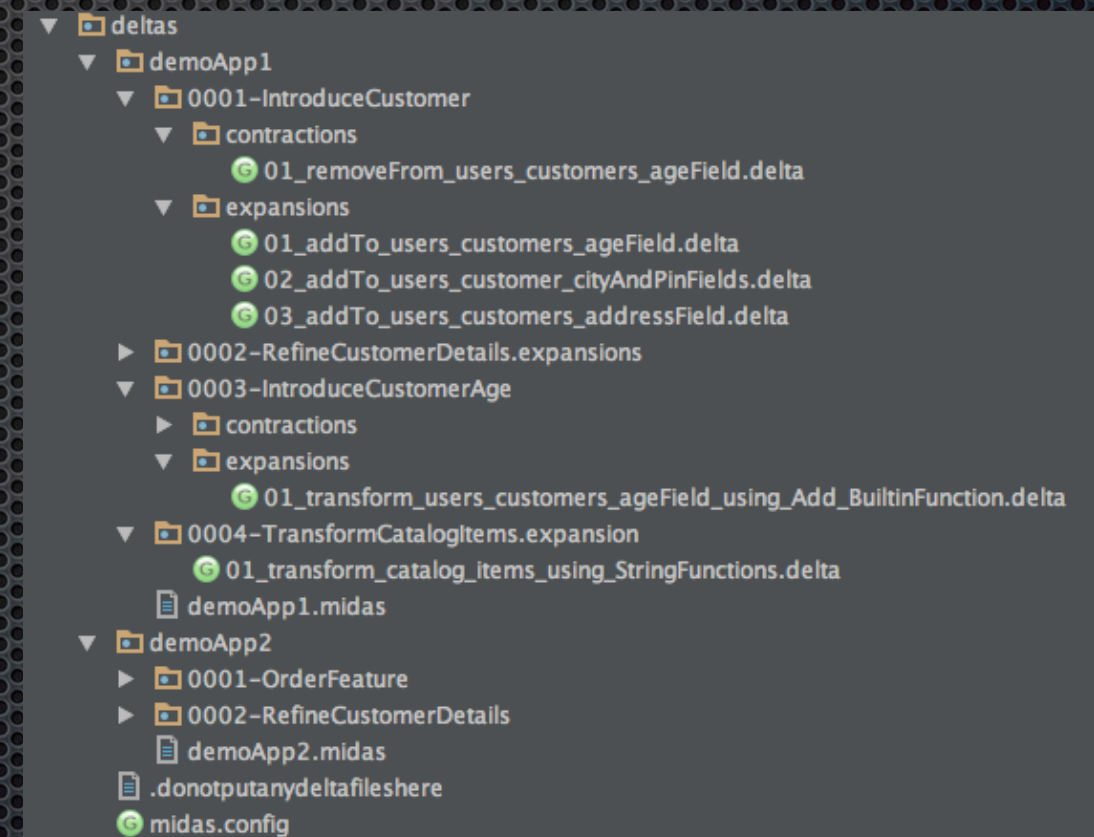
- ✦ From the App perspective
 - ✦ Midas is Agnostic of Language specific drivers and versions within those languages
 - ✦ Works with versions of Ruby, Python, C# and Java drivers.
- ✦ From MongoDB perspective
 - ✦ Midas is Agnostic of the MongoDB configurations
 - ✦ Works with Standalone, Replica Sets, Sharded environments.

Deep Dive



Delta scripts

- ✦ Group ChangeSet as Directories
- ✦ Group deltas as Expansion and/or Contraction deltas as Directories within ChangeSet
- ✦ Write delta scripts using the `.delta` extension
- ✦ Midas relies on the order specified by you
- ✦ Ordering info is embedded by you within the change set directory name and delta files



Delta Scripts - Convention

- ✦ Must begin change set directories with a number and subsequent change set directories are numbered in ascending order.
 - ✦ Example: use `<changesetNumber>-<featureName>` as convention for naming change set directory.
- ✦ Within a change set, expansion and contraction scripts are grouped by folders - `expansions` and `contractions`.
- ✦ Always begin delta scripts with a number and subsequent delta scripts are numbered in ascending order.
 - ✦ Example: use `<changeNumber>_<WhatTheChangeIs>.delta` as convention for naming delta files

Sample Delta Script

- Each Delta script is written using a DSL
- Very close to MongoDB lingo, virtually no learning curve.

```
use users
db.customers.remove(`["address.line1"]`)
db.customers.merge(["lname", "fname"], " ", "name")

use transactions
db.orders.add(`dispatch : { status: 'NOT DELIVERED' }`)
```


Agile App Delivery & DevOps

- ✦ Inject Midas into Architecture
 - ✦ Start or Middle of project
- ✦ Supports Development of Application in small-steps
 - ✦ Add Application ChangeSets/Deltas on-the-fly
- ✦ Copes with Load
 - ✦ Add/Remove Application Nodes on-the-fly
- ✦ Supports Multiple Applications
 - ✦ Add/Remove Application(s) on-the-fly

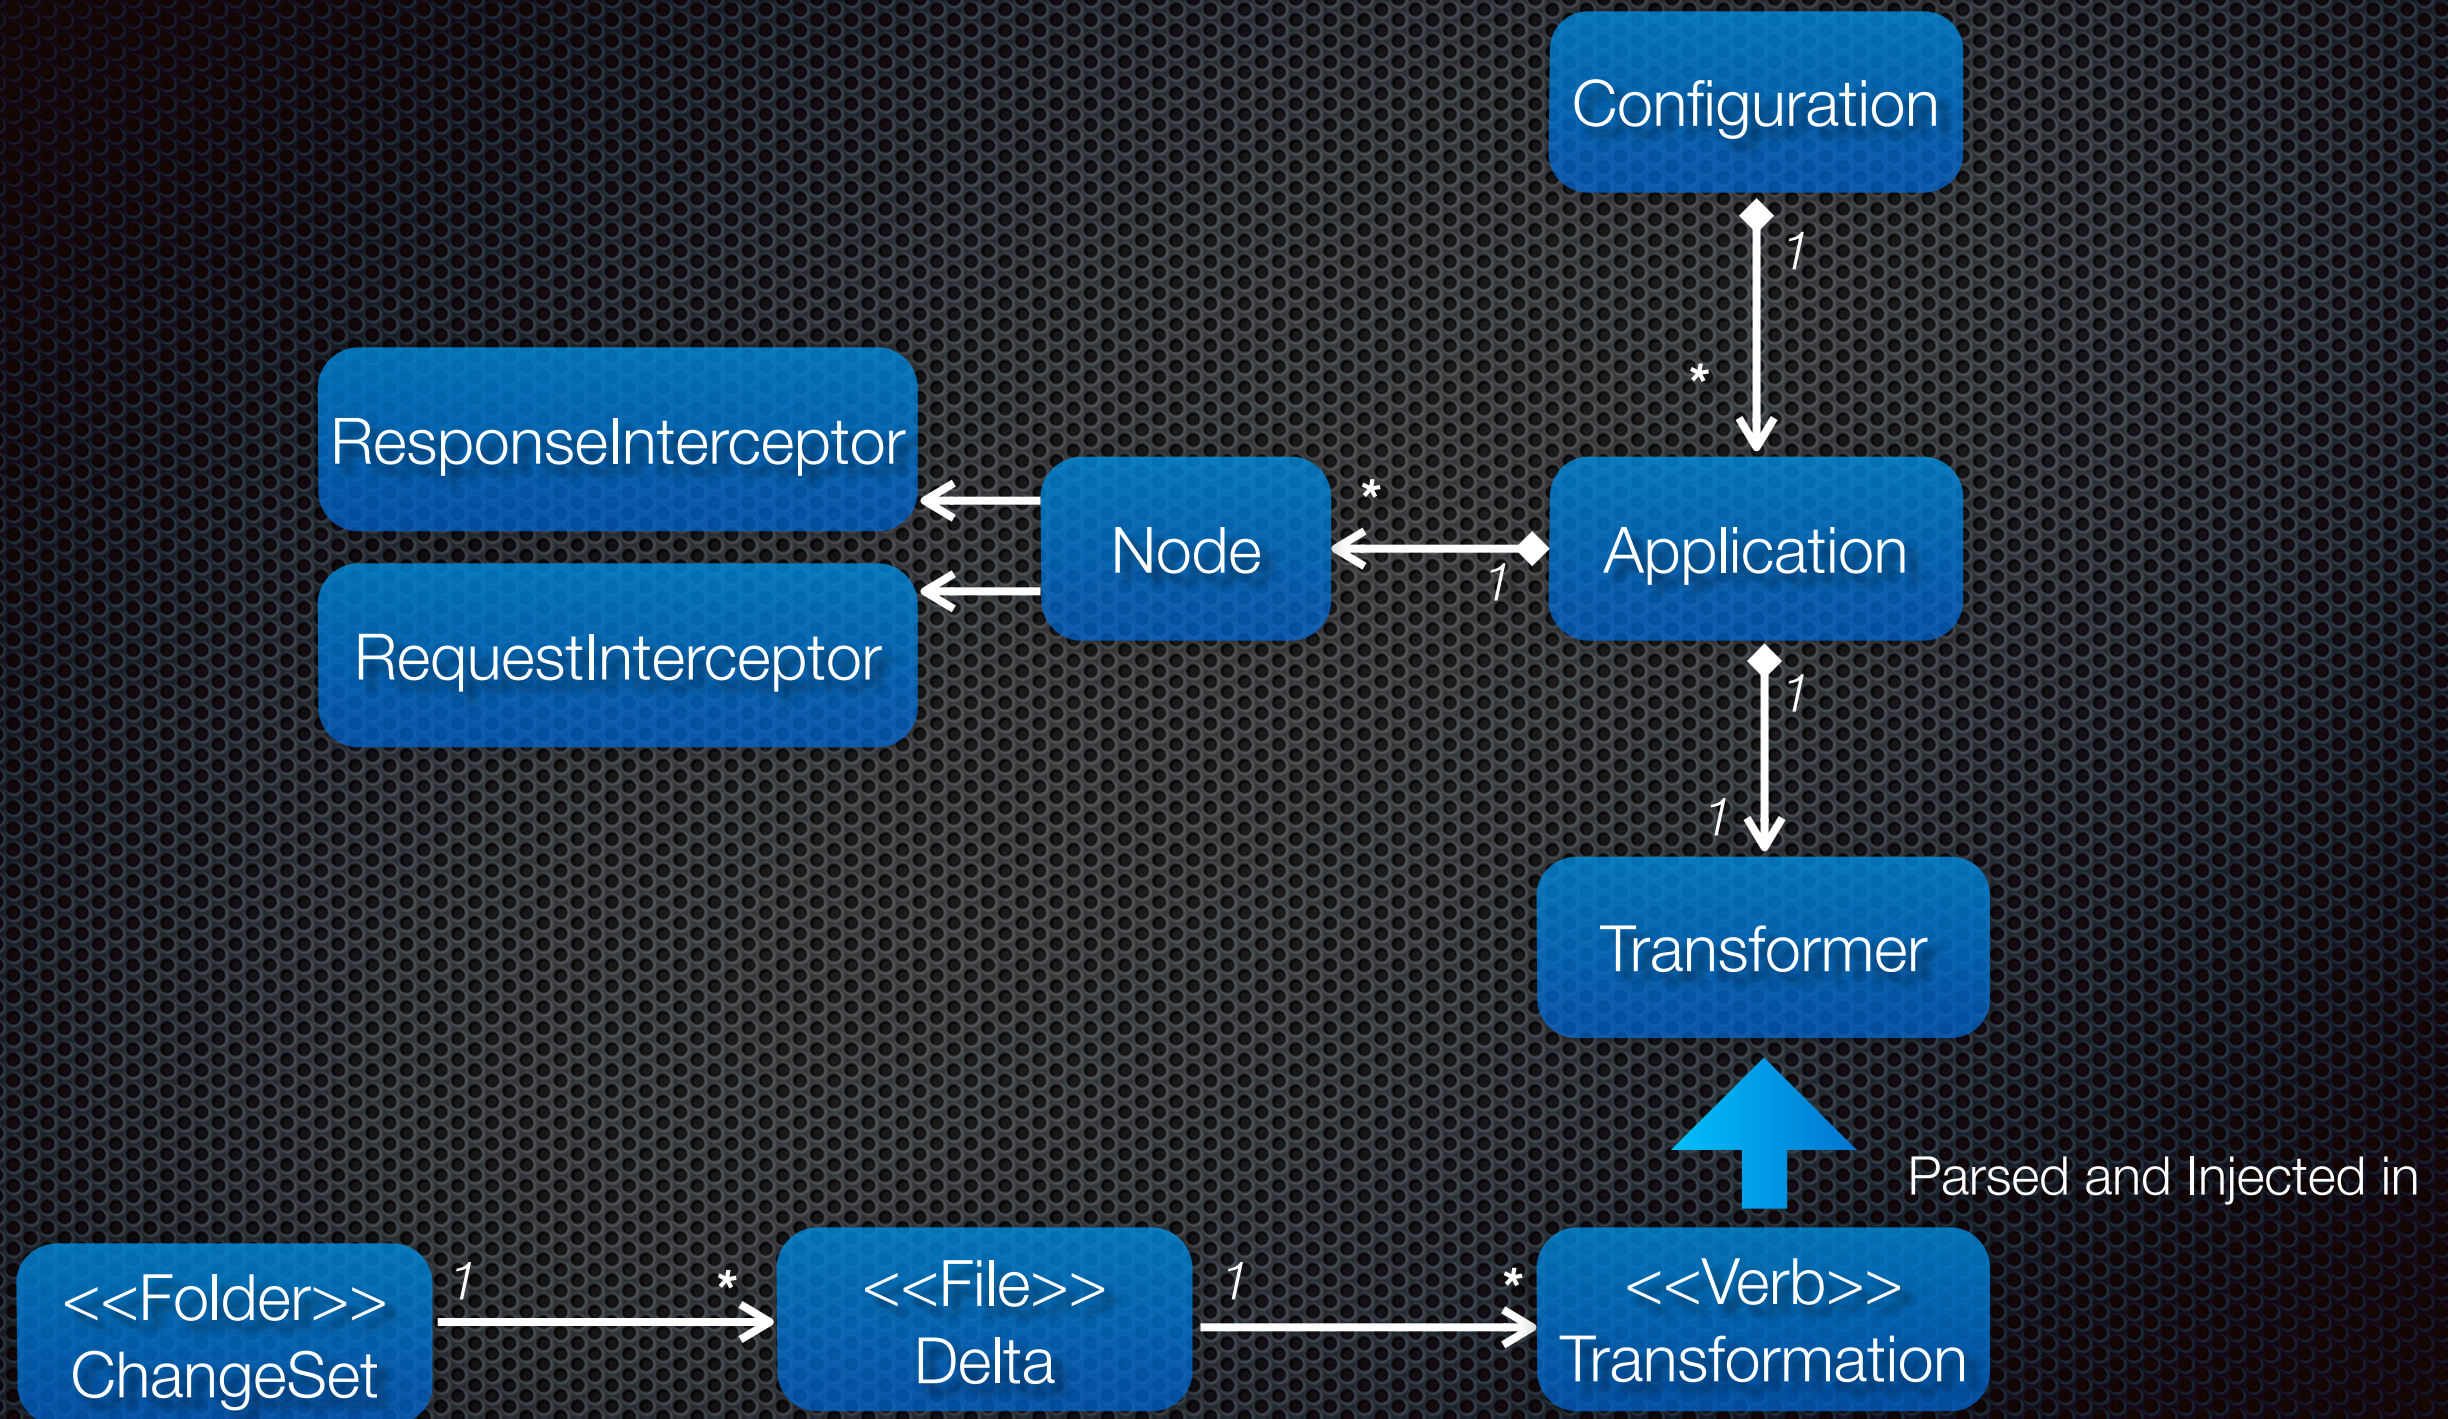
No Leaky Abstractions

- ✦ Does not expect the Domain Model to be aware of versioning.
 - ✦ Allows developers to focus on the domain while freeing them from versioning concerns.
 - ✦ If you wish to take charge, Midas will not come in your way.
- ✦ Midas maintains versioning information within the document itself.
 - ✦ `_expansionVersion` and `_contractionVersion` fields are part of the documents “touched” by Midas.
 - ✦ Updates them during request and response.

Caveats

- Never ever change a delta that has been applied to production.
 - Always move forward in time.
 - Reverse a change by a counter-change.
- Force-update migration on documents that are not
 - expanded by App demand, and then proceed to contraction.
 - contracted by App demand, and then proceed to next App-upgrade cycle.
- Currently Midas supports Zero Downtime Configuration - 1 (see later slides), others will be supported in subsequent releases.

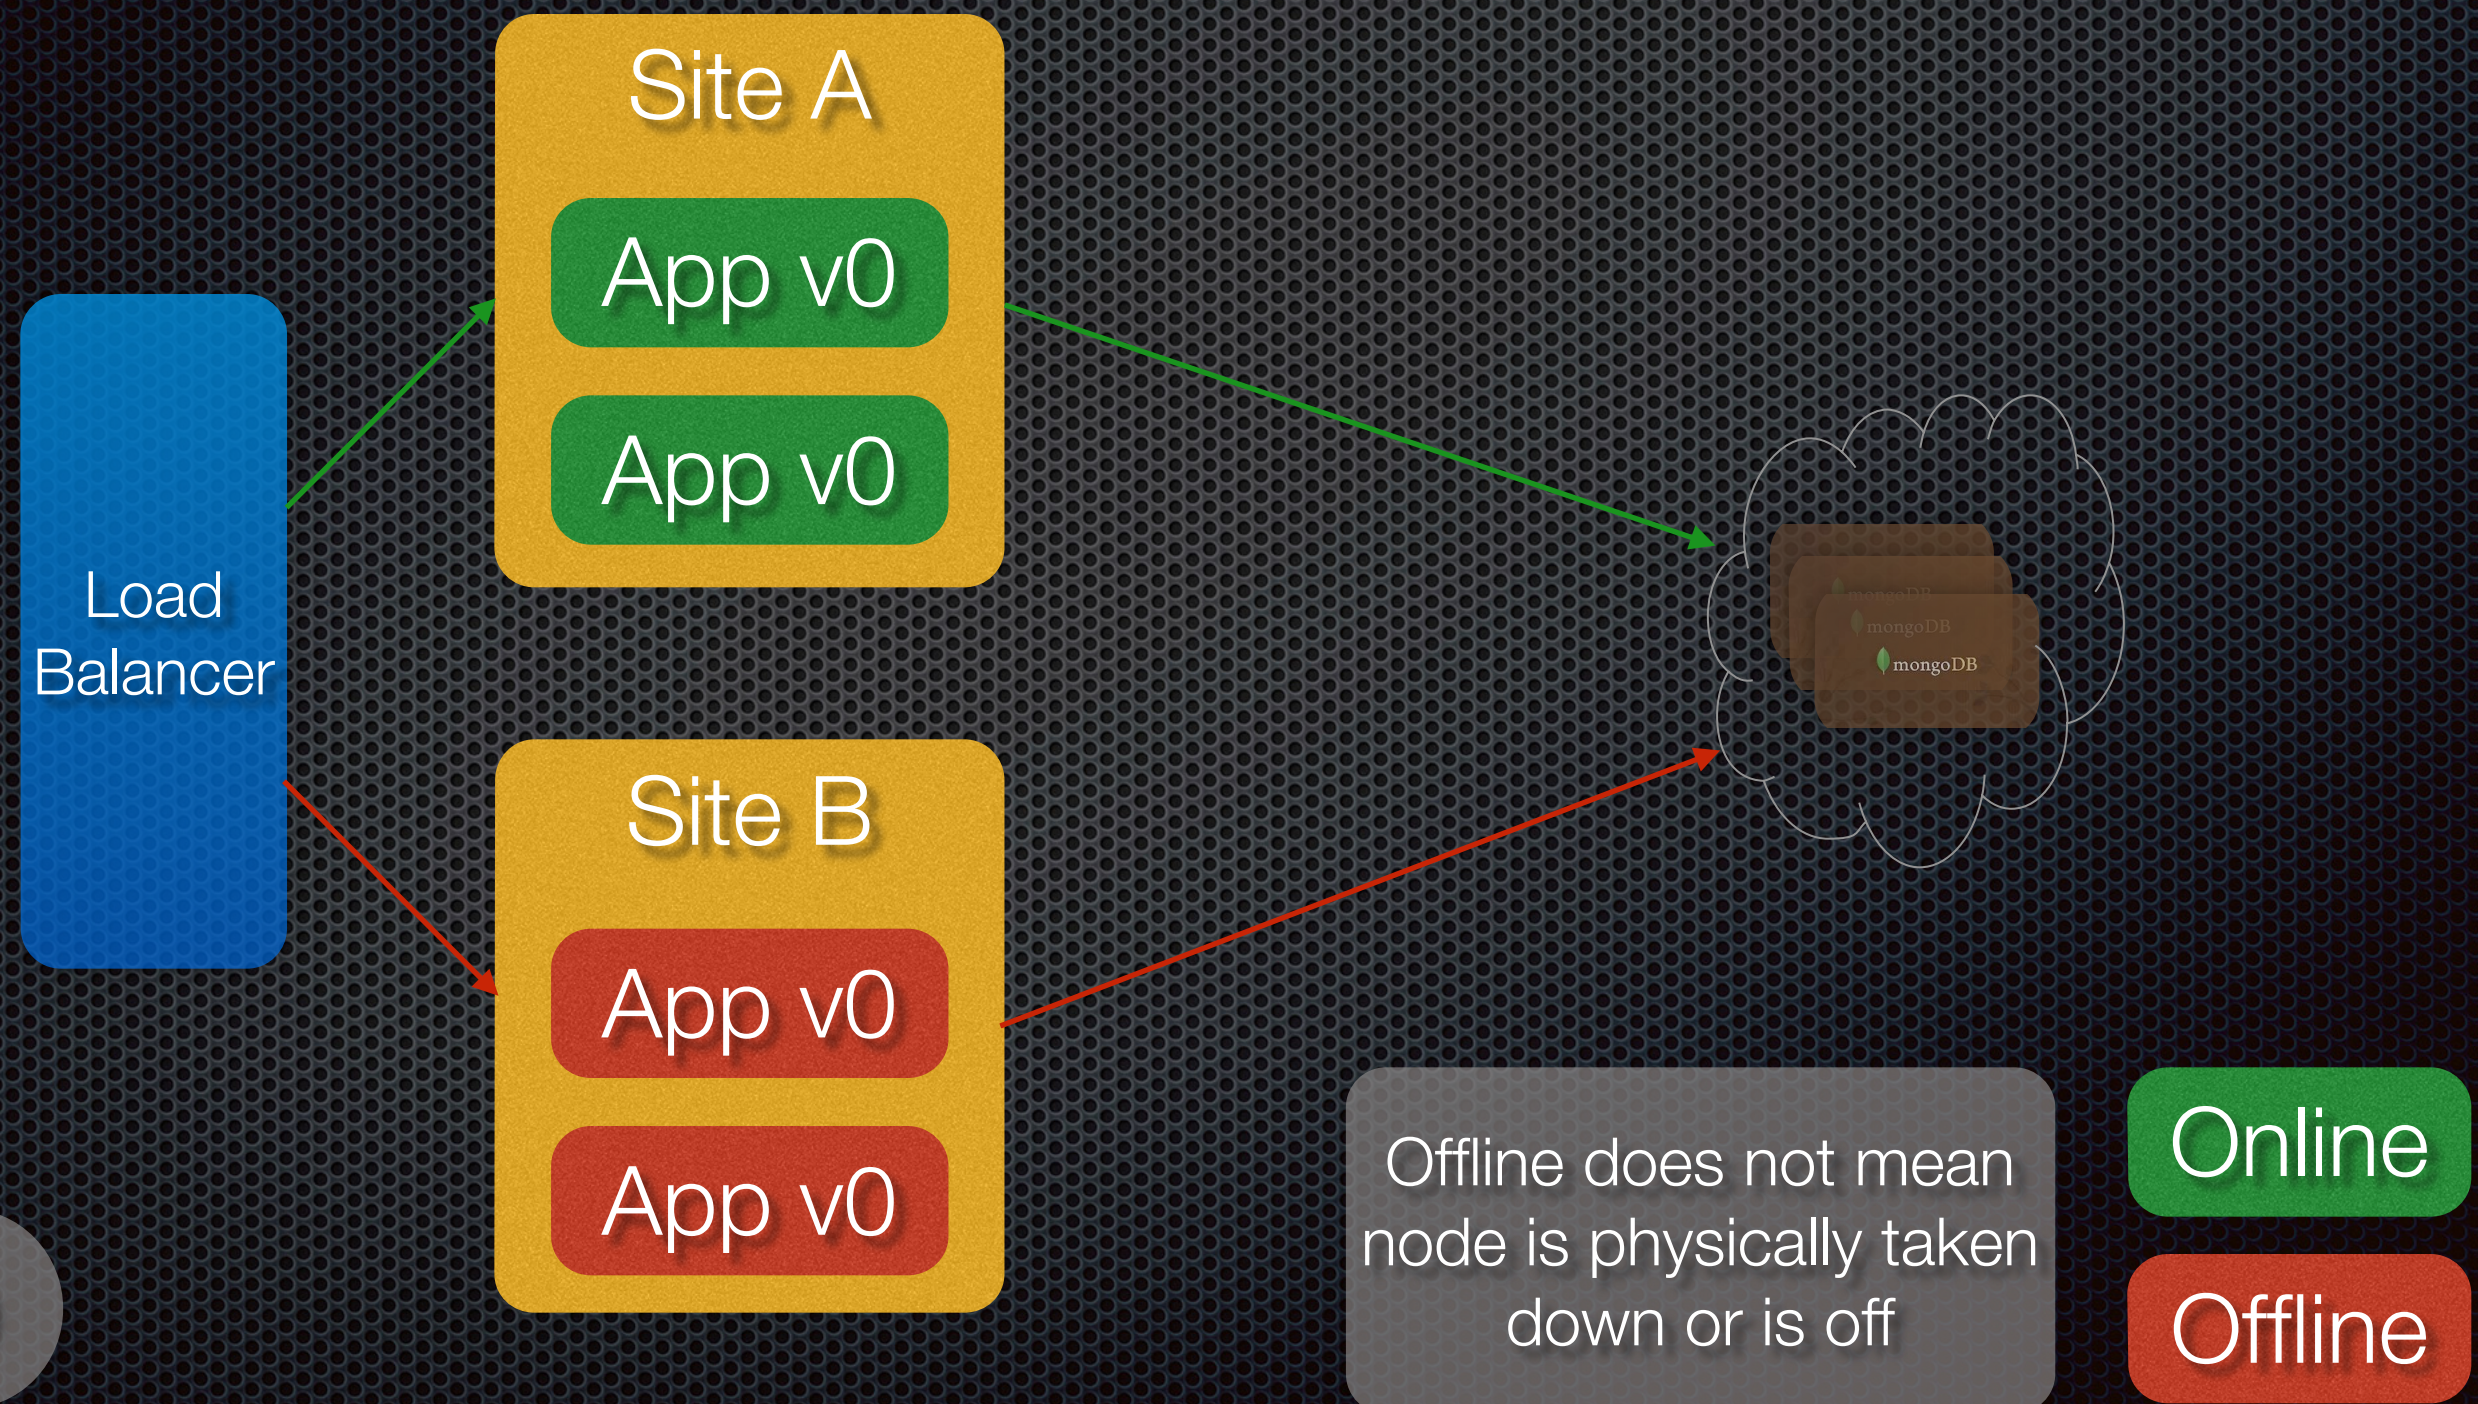
Midas Domain Model



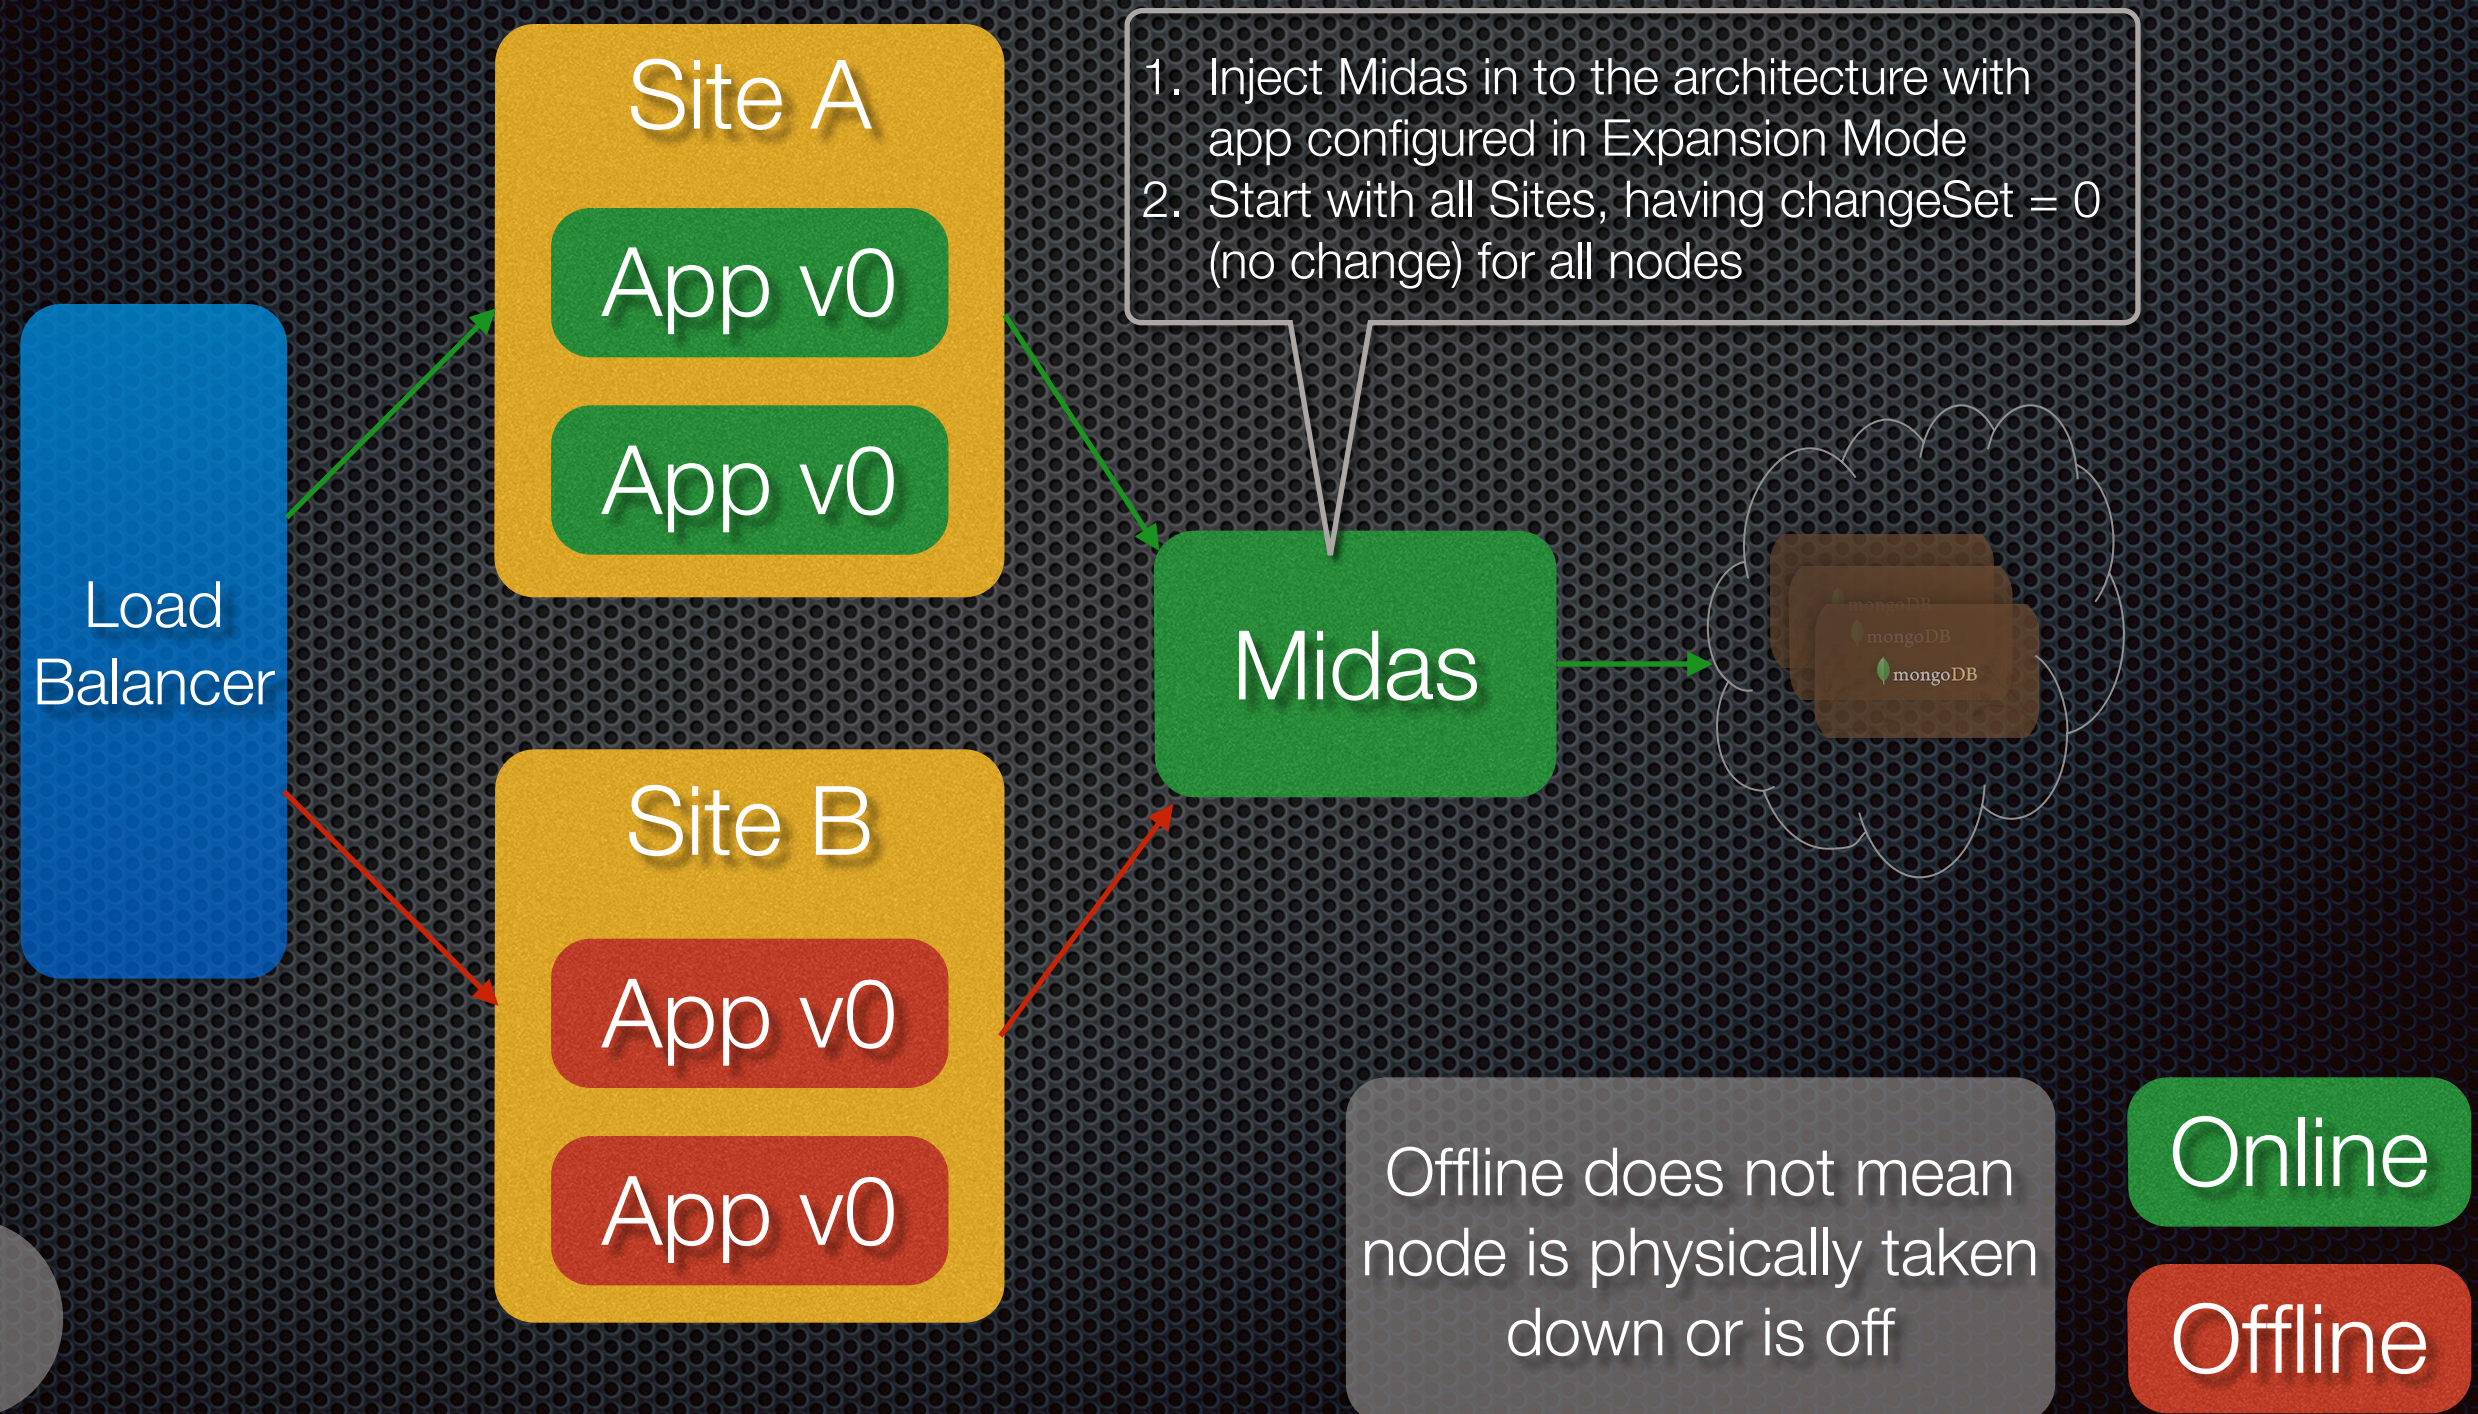
Zero-Downtime Deployment Configuration - 1

Injecting and Using Midas

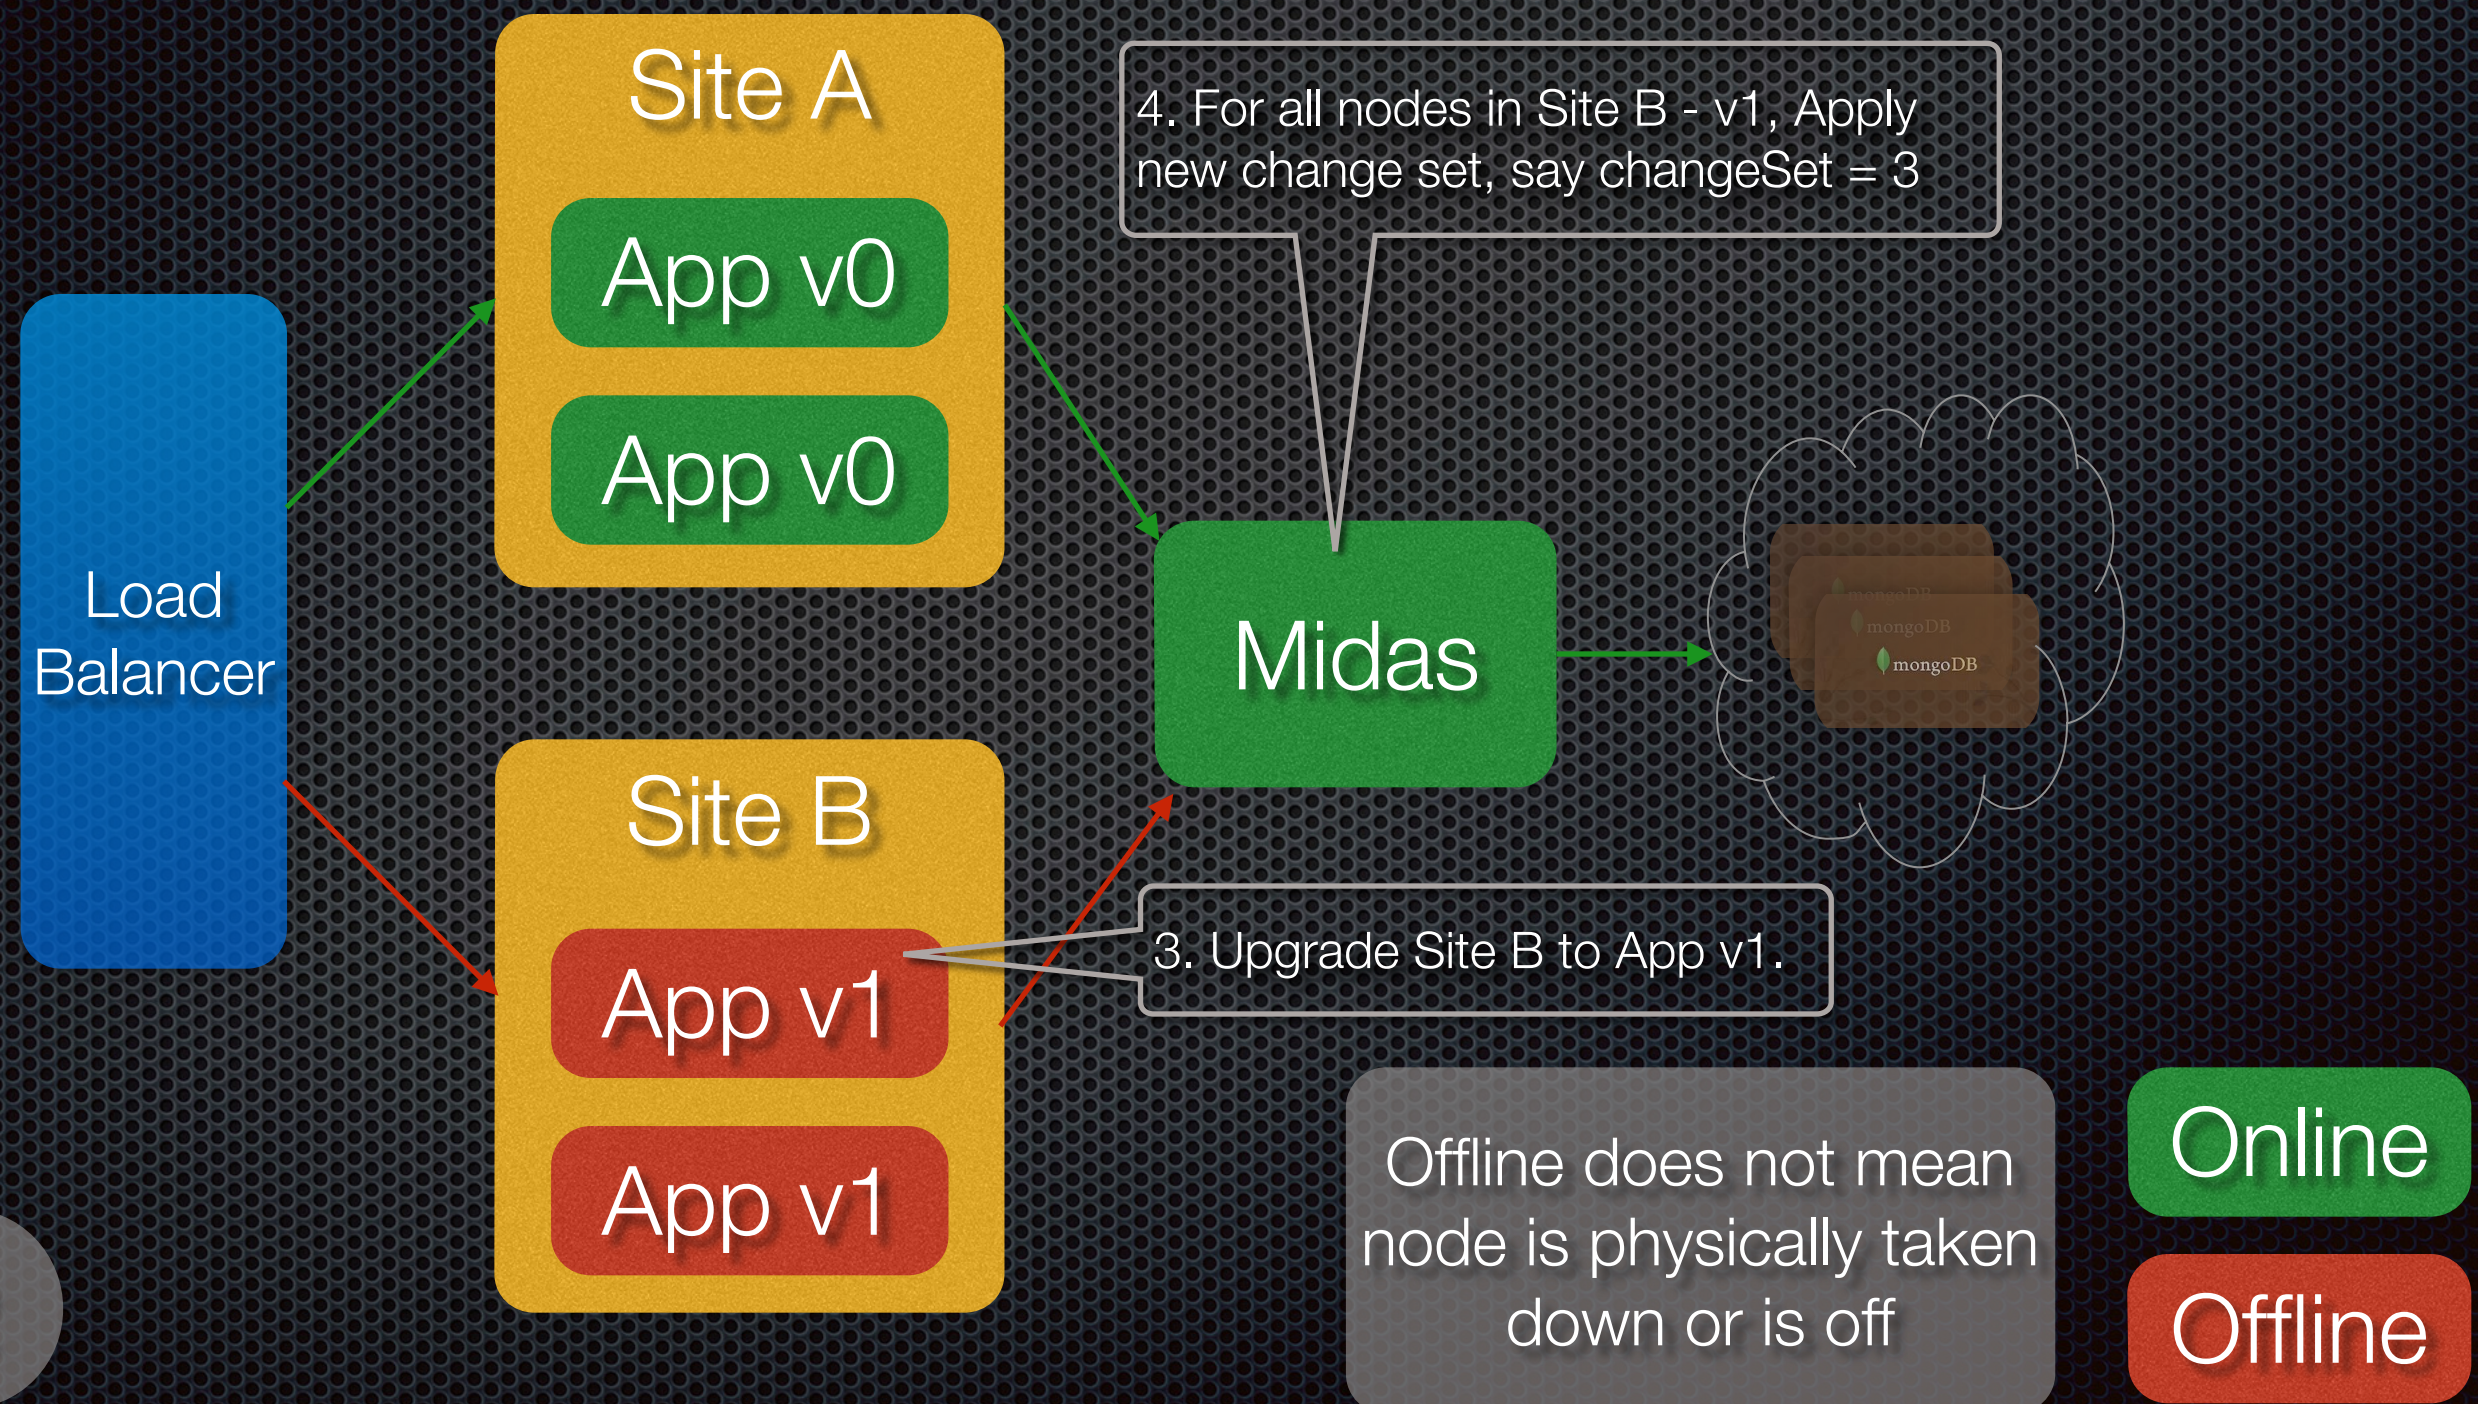
Deployment config - 1



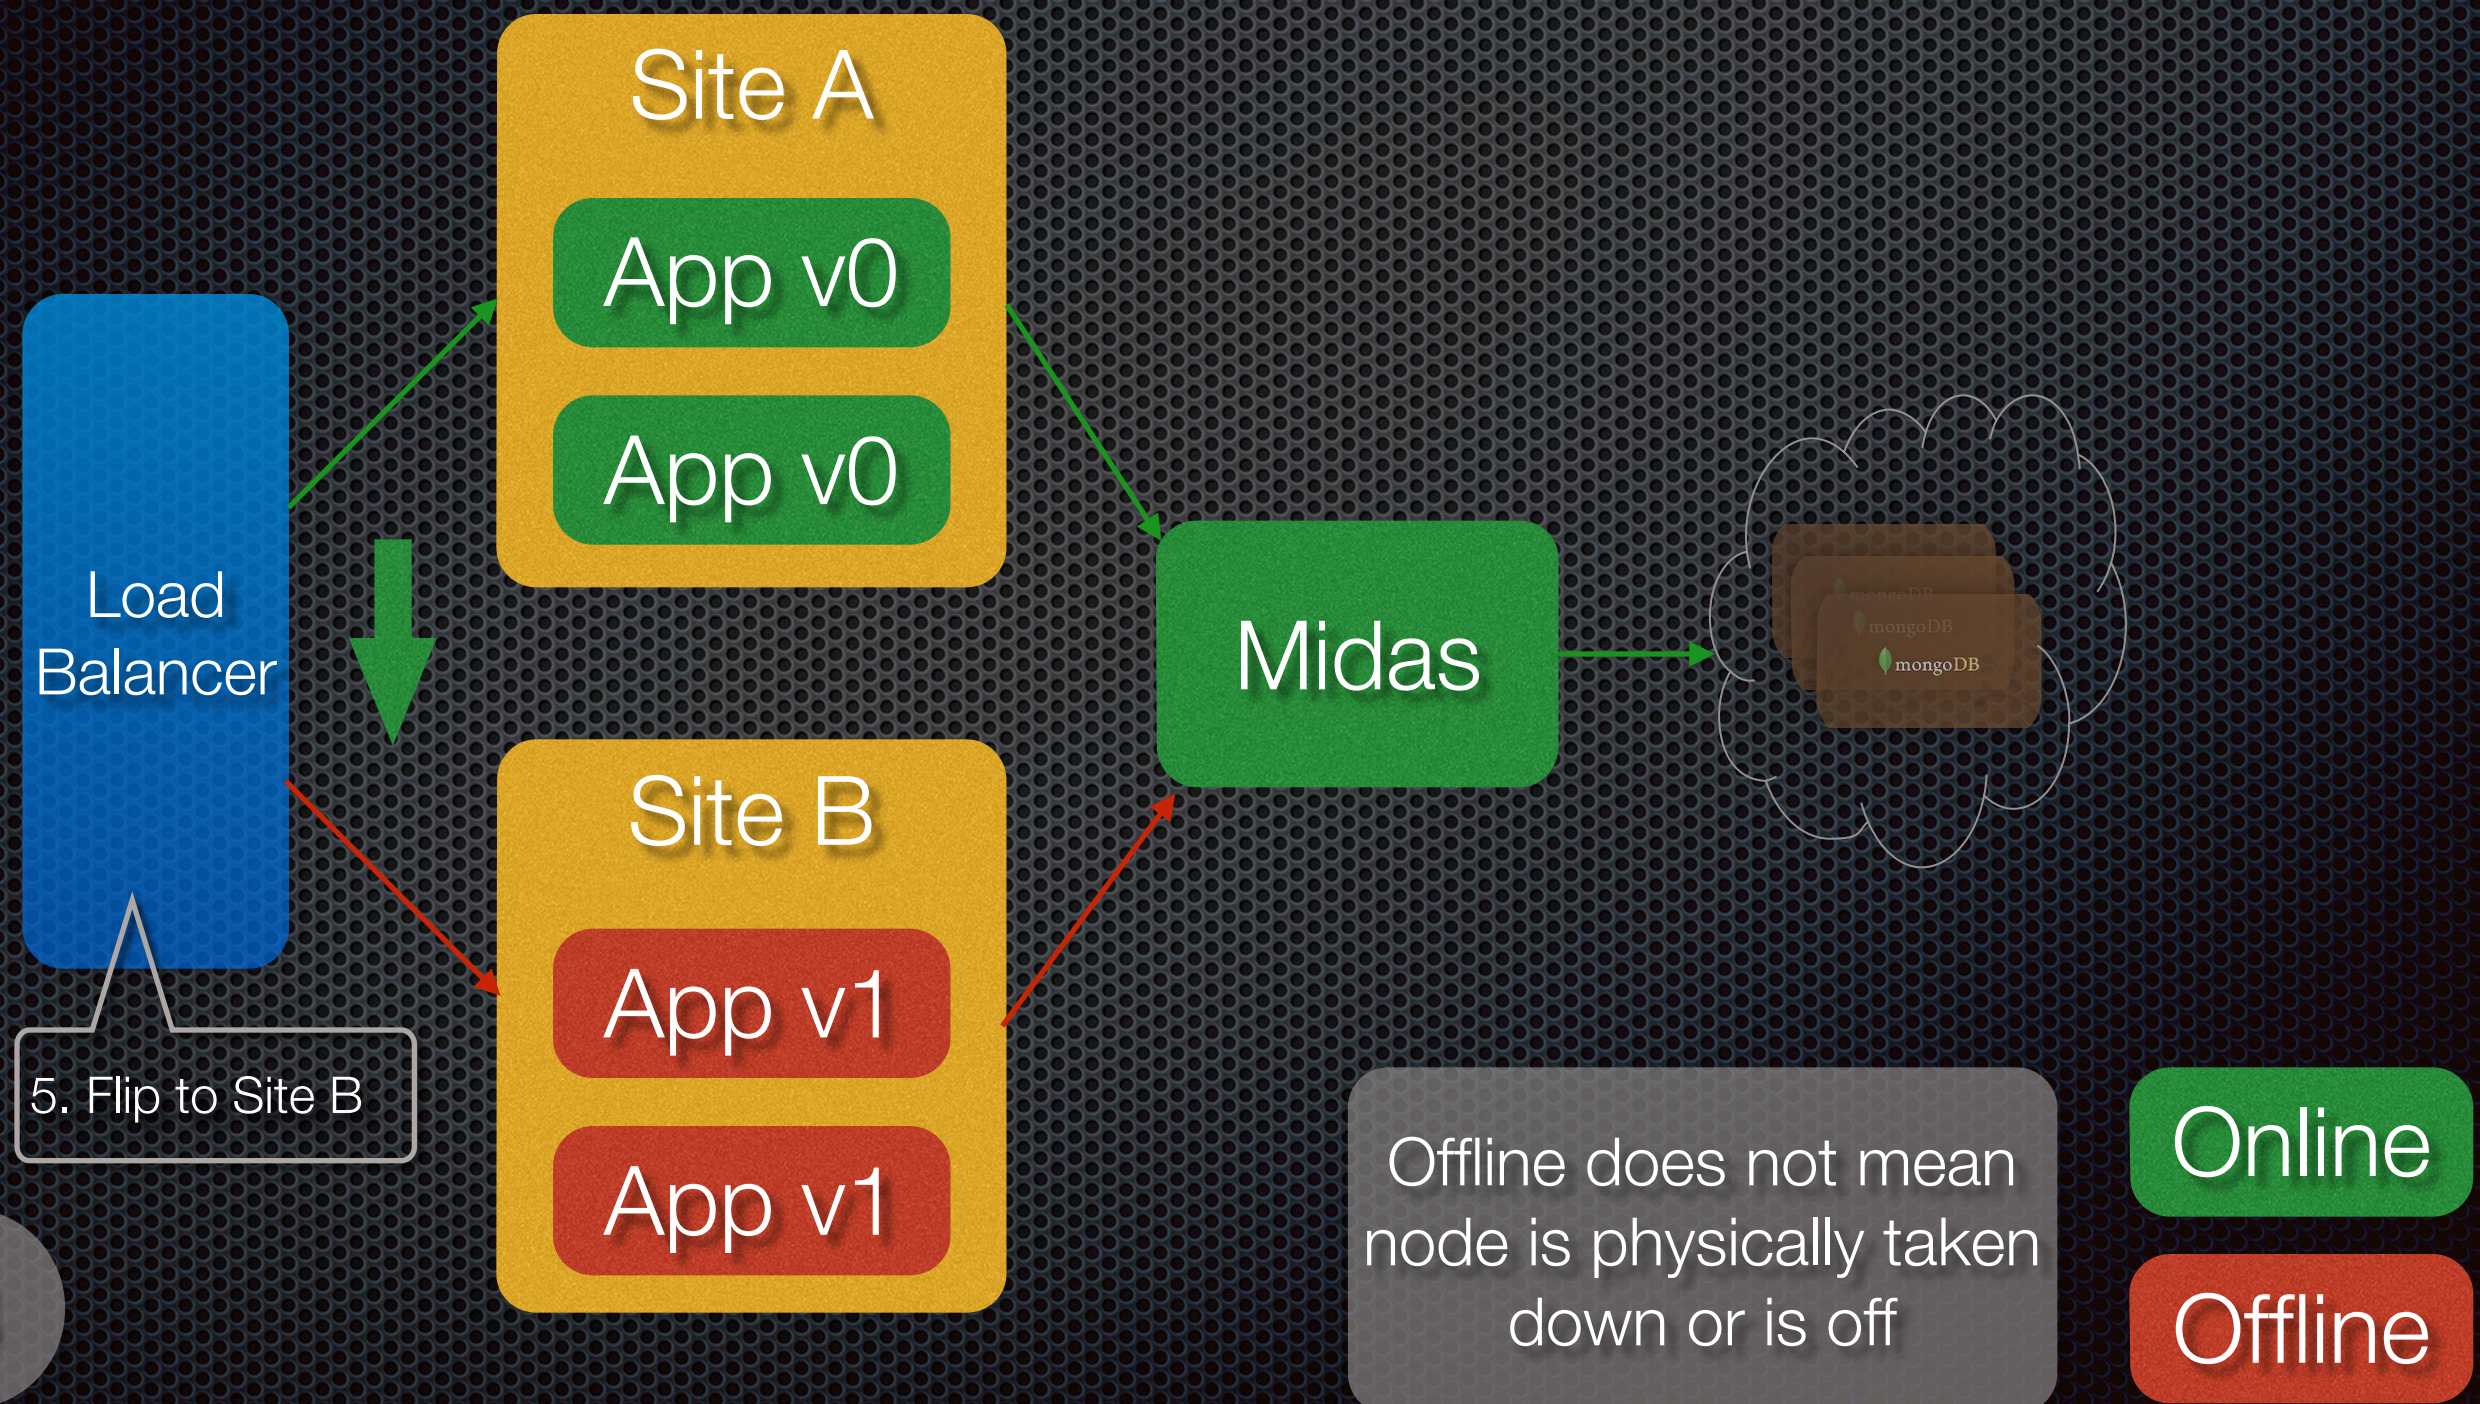
Deployment config - 1



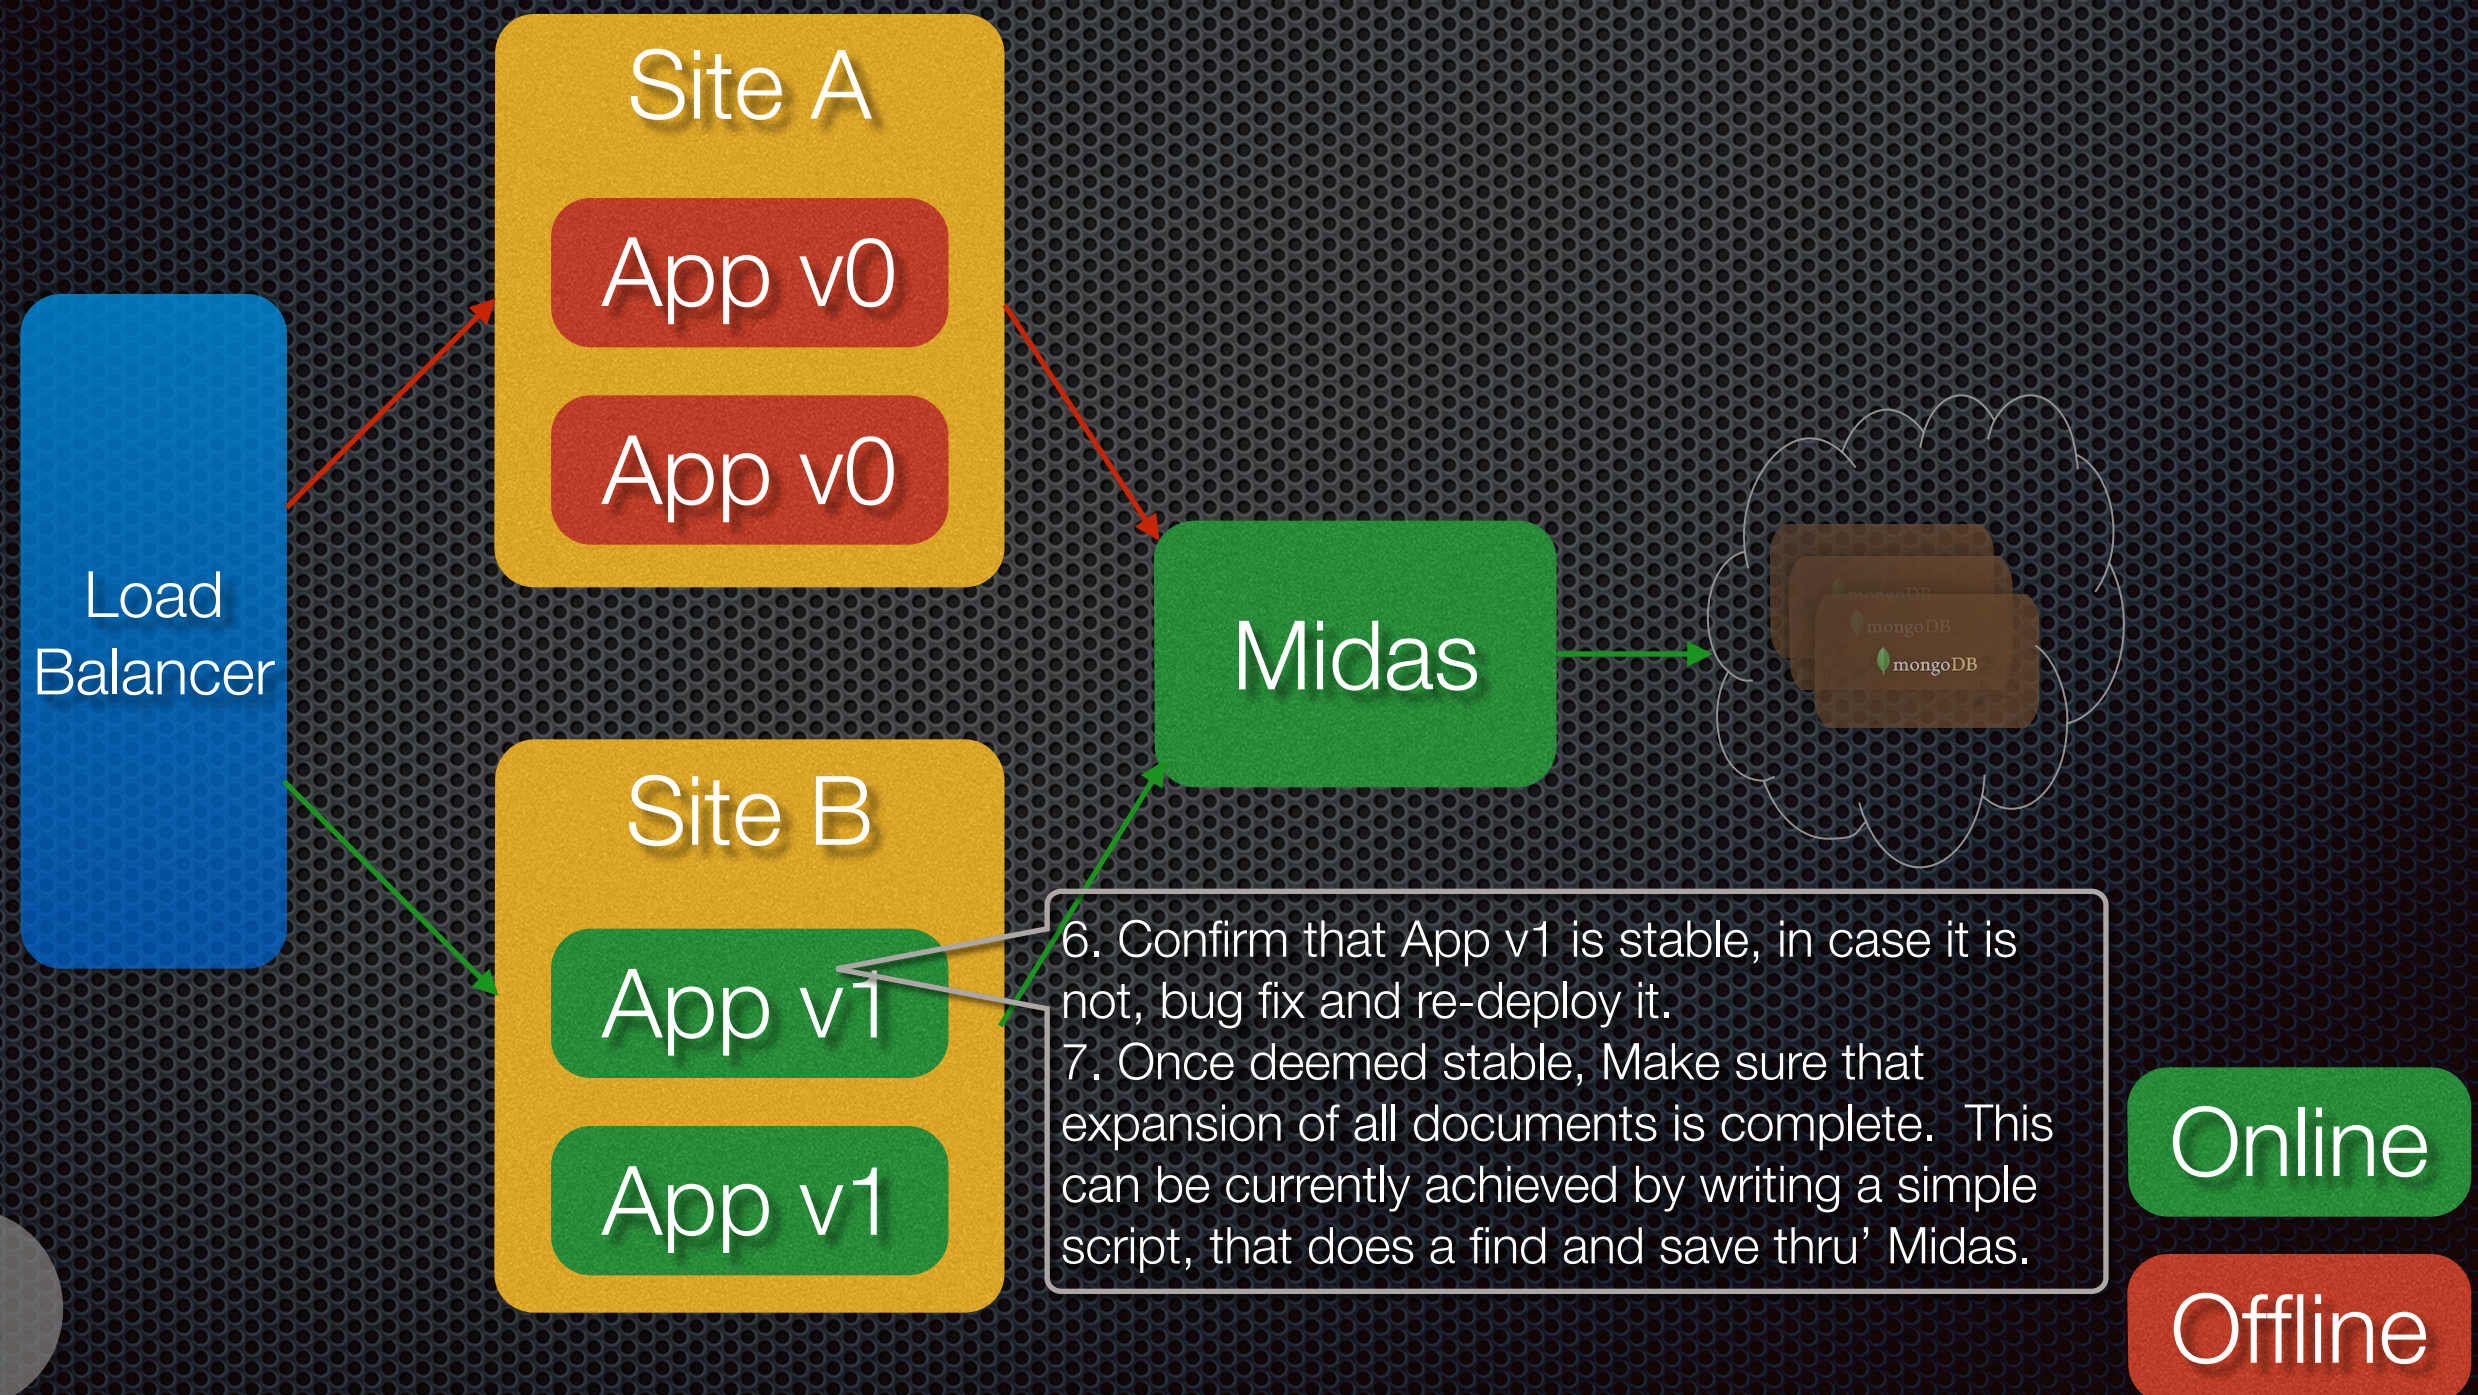
Deployment config - 1



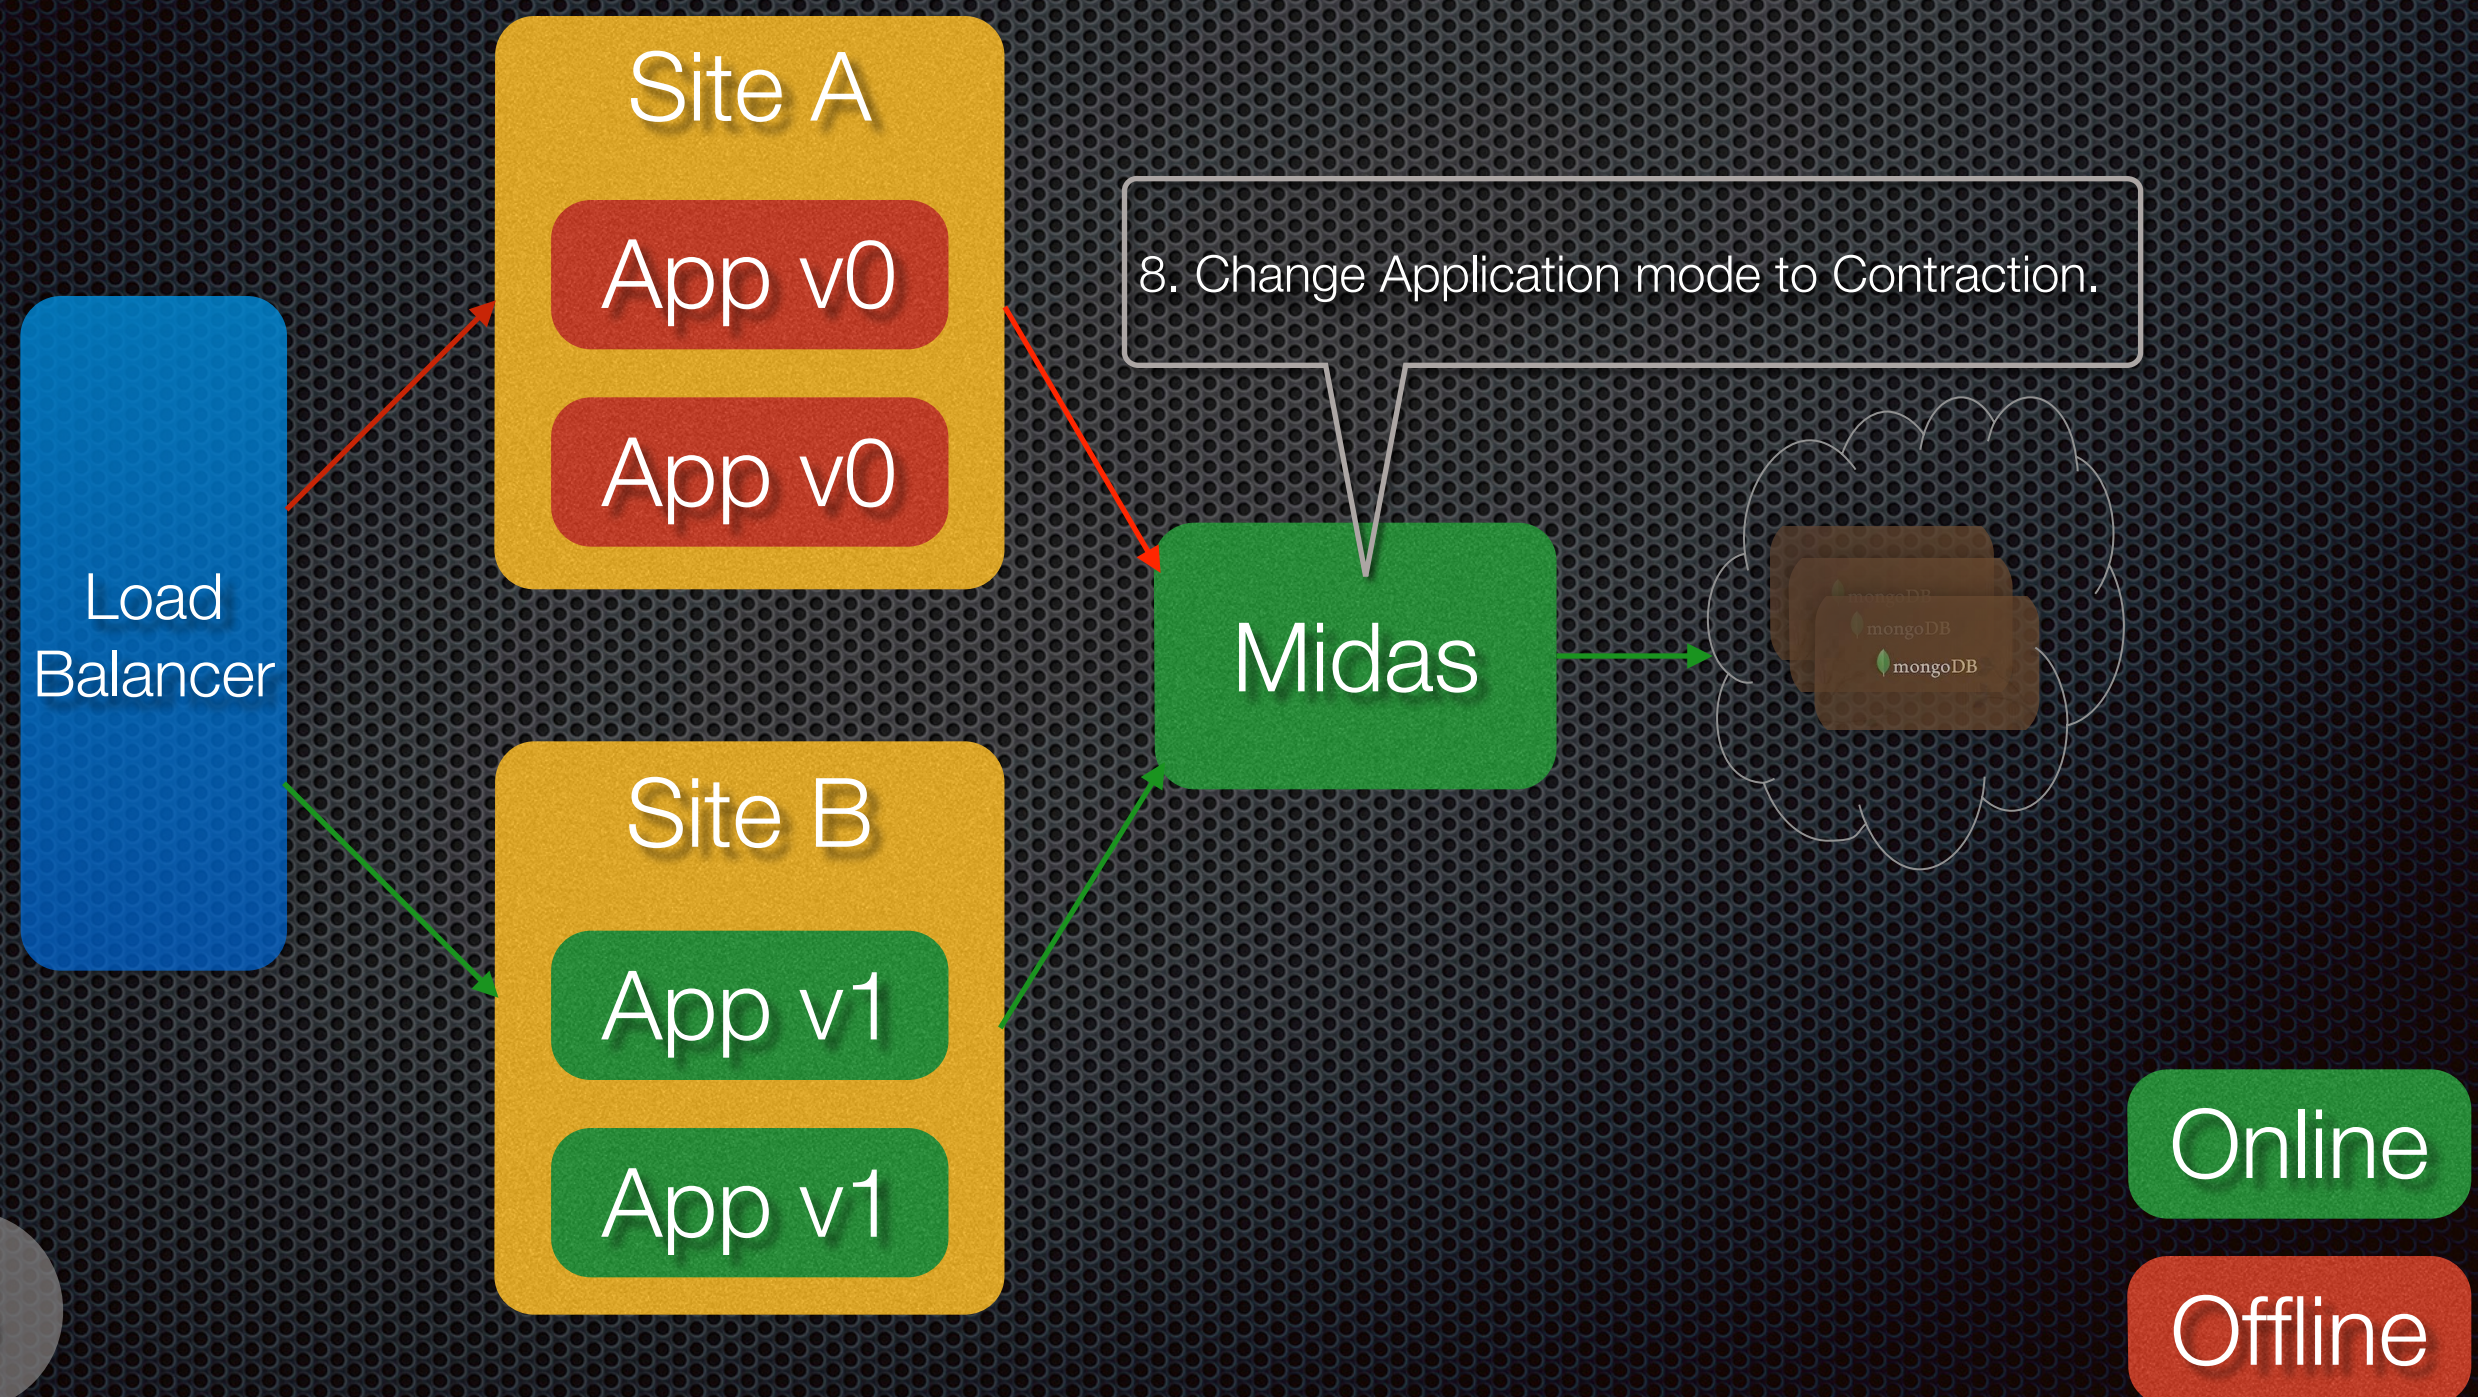
Deployment config - 1



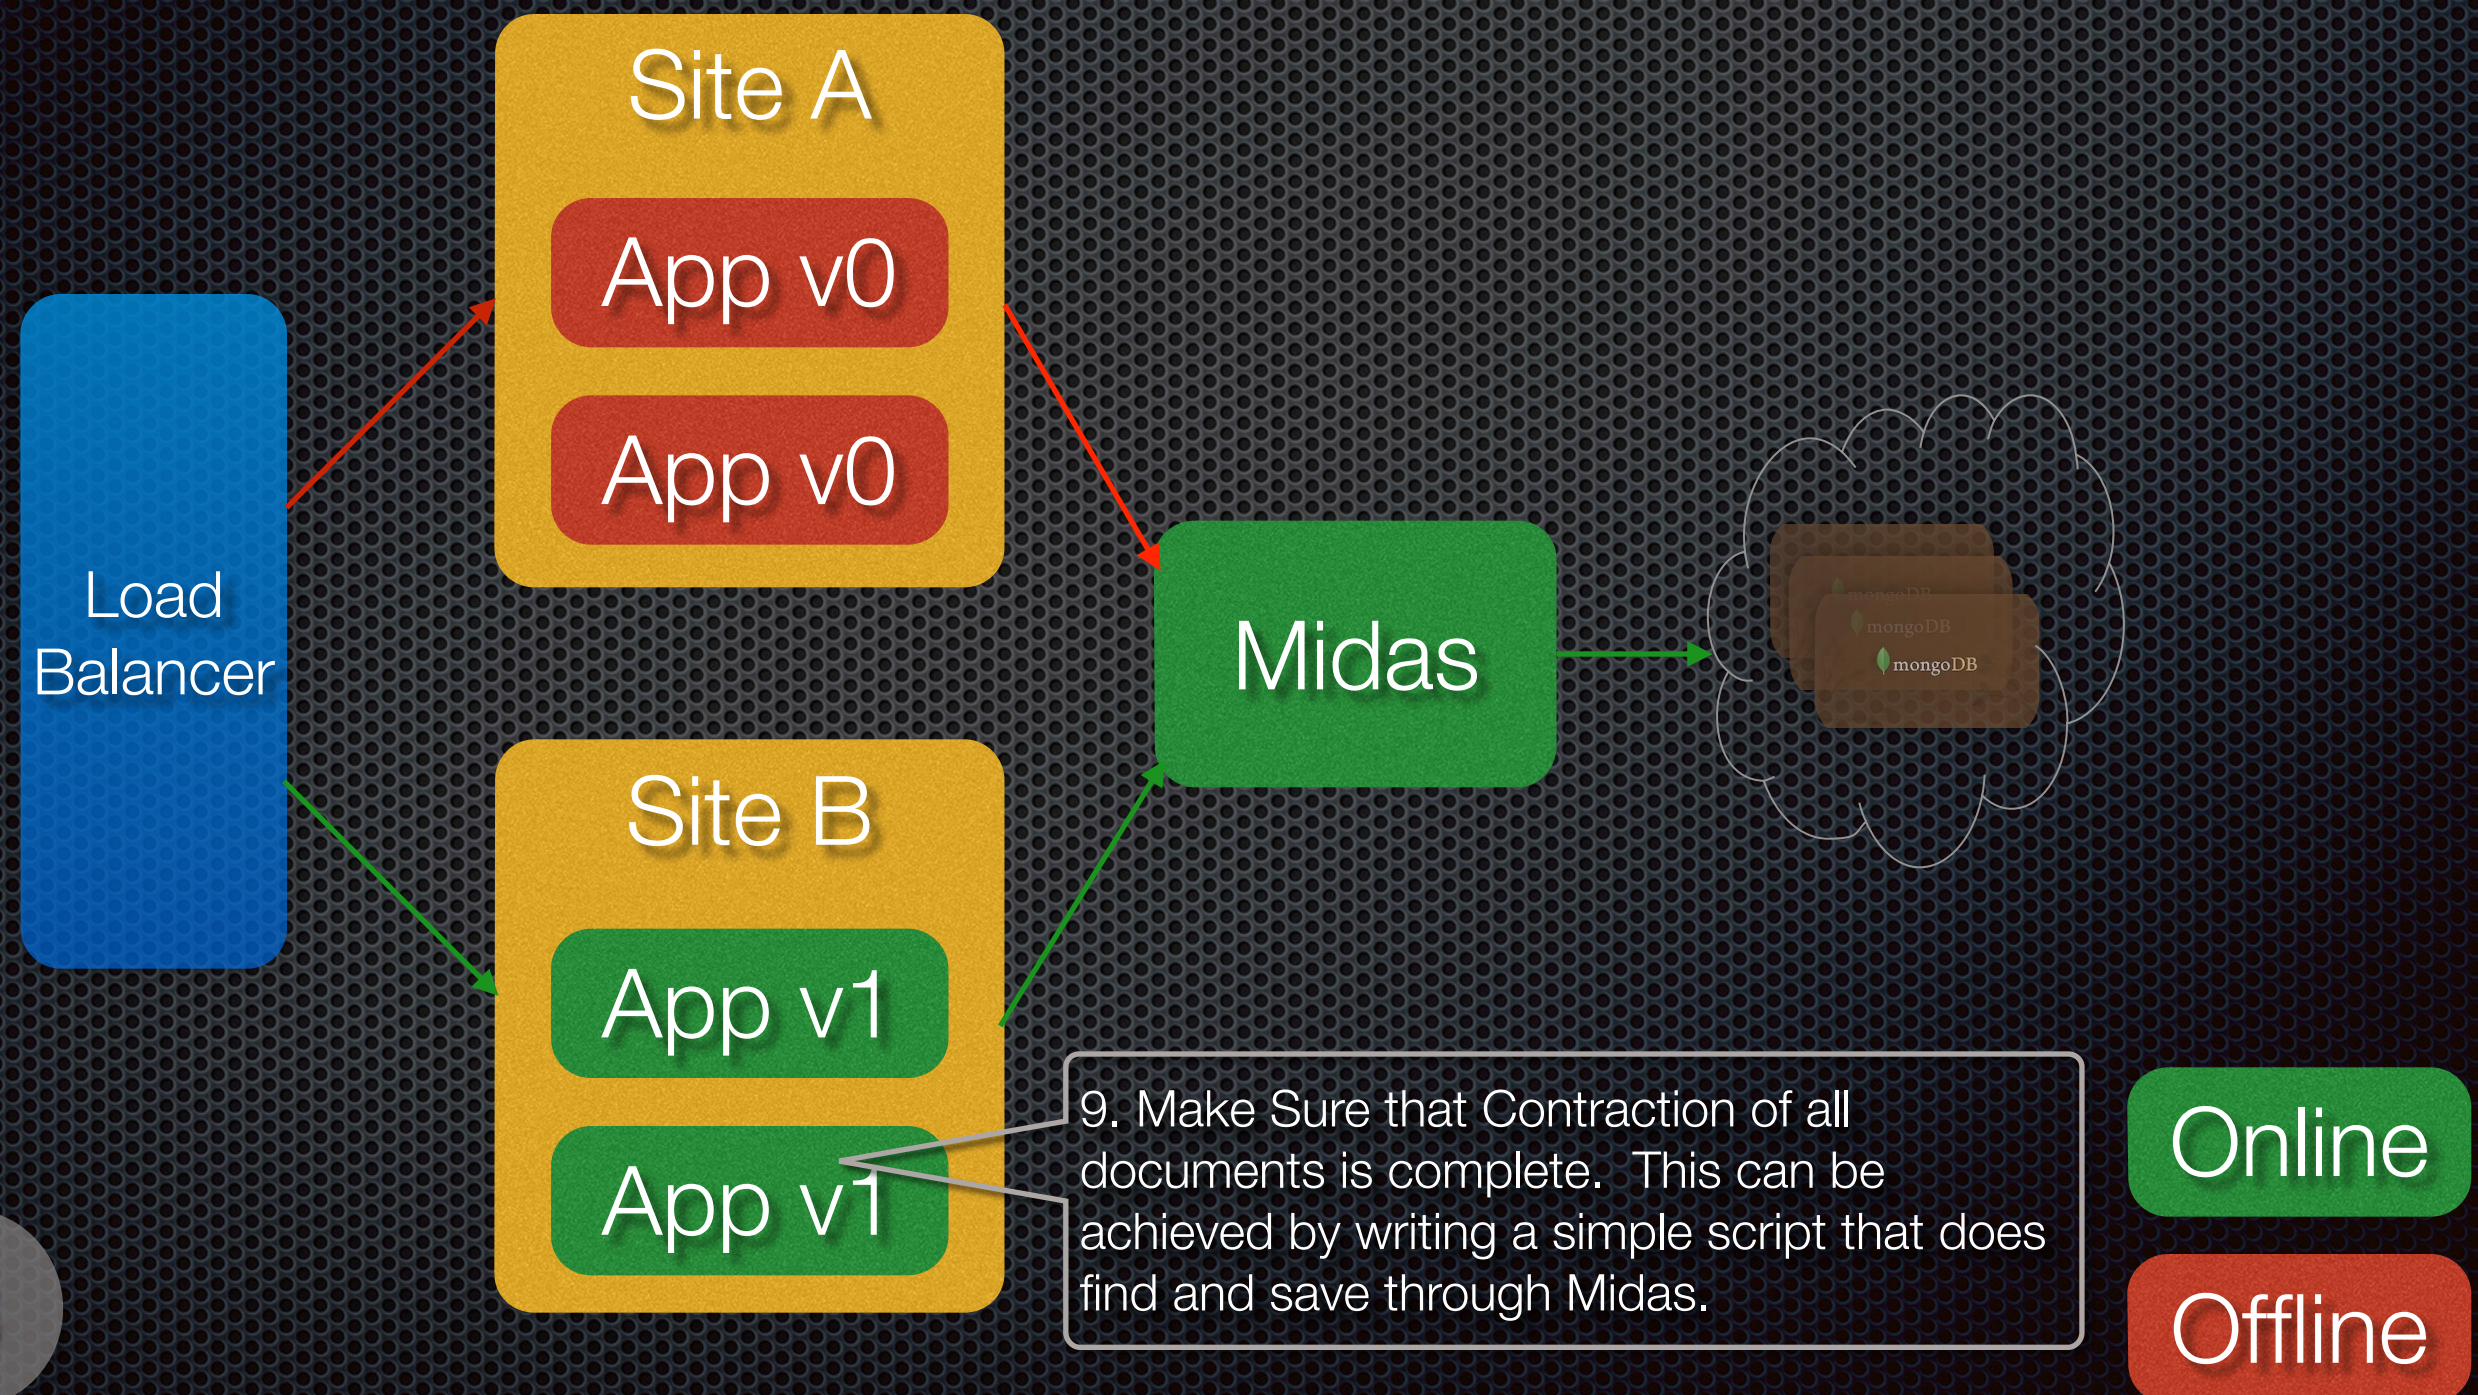
Deployment config - 1



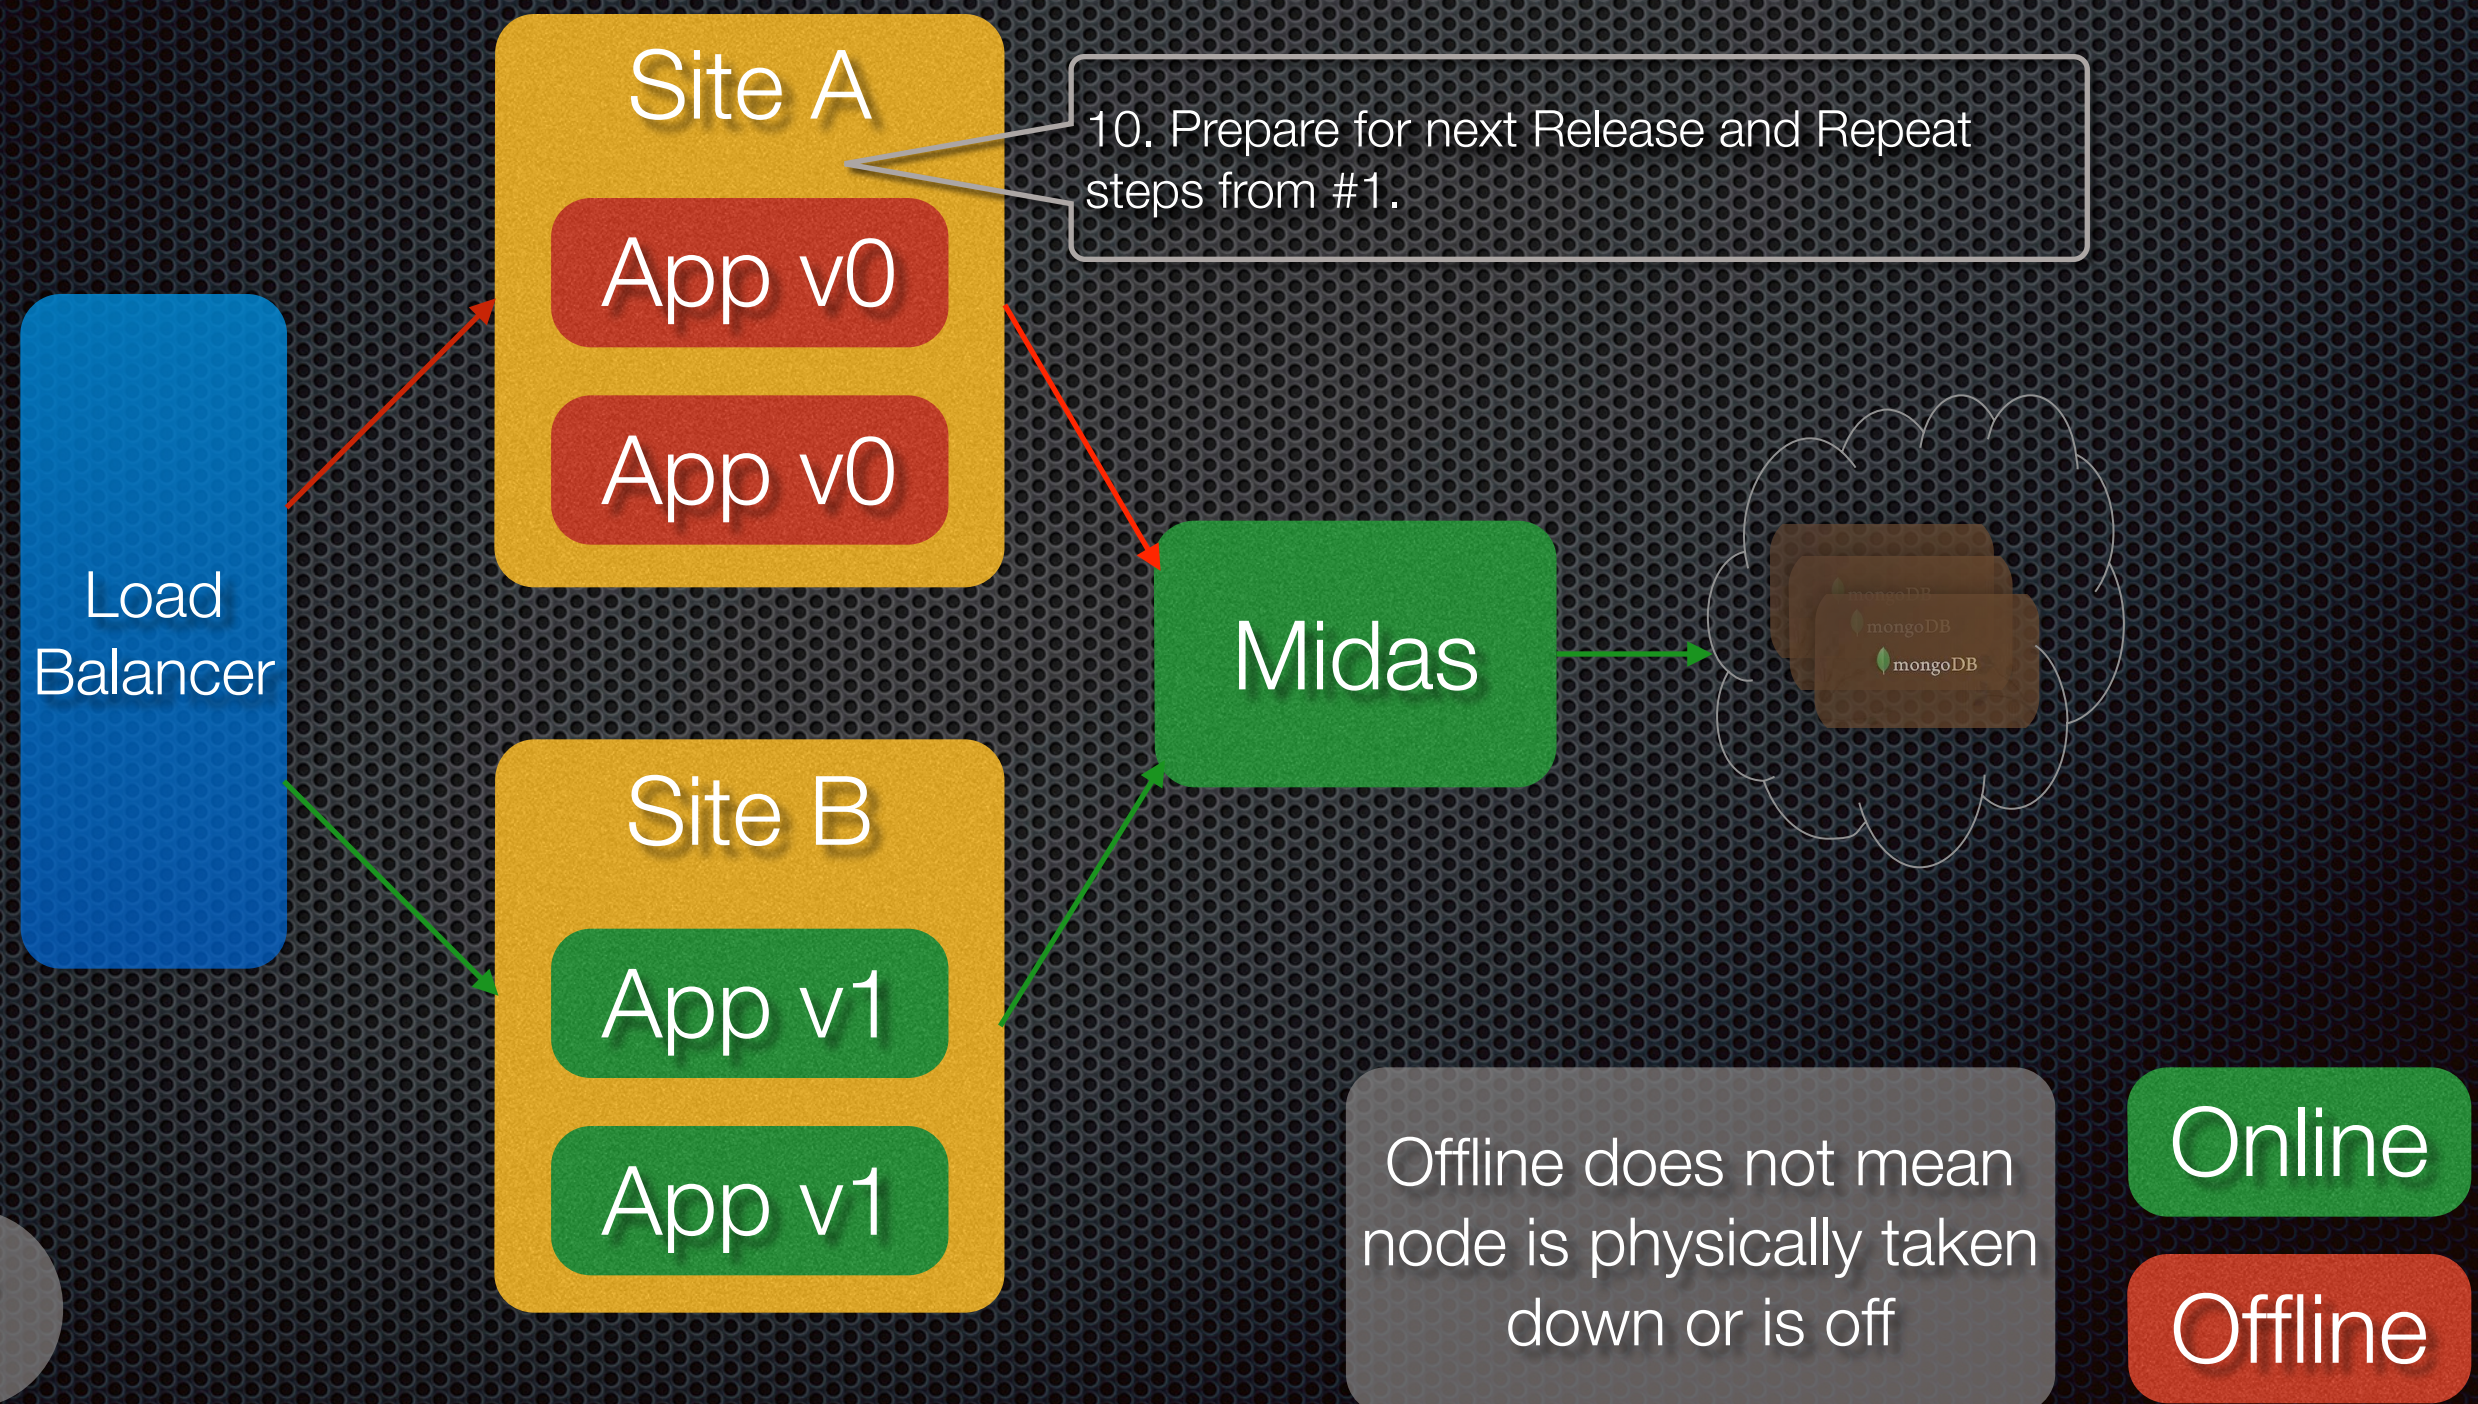
Deployment config - 1



Deployment config - 1



Deployment config - 1



Zero-Downtime Deployment Configuration - 1




Node Removal at Runtime

Say, due to some problems, you want to remove Node0 from service:

1. Remove Node0 from LB.

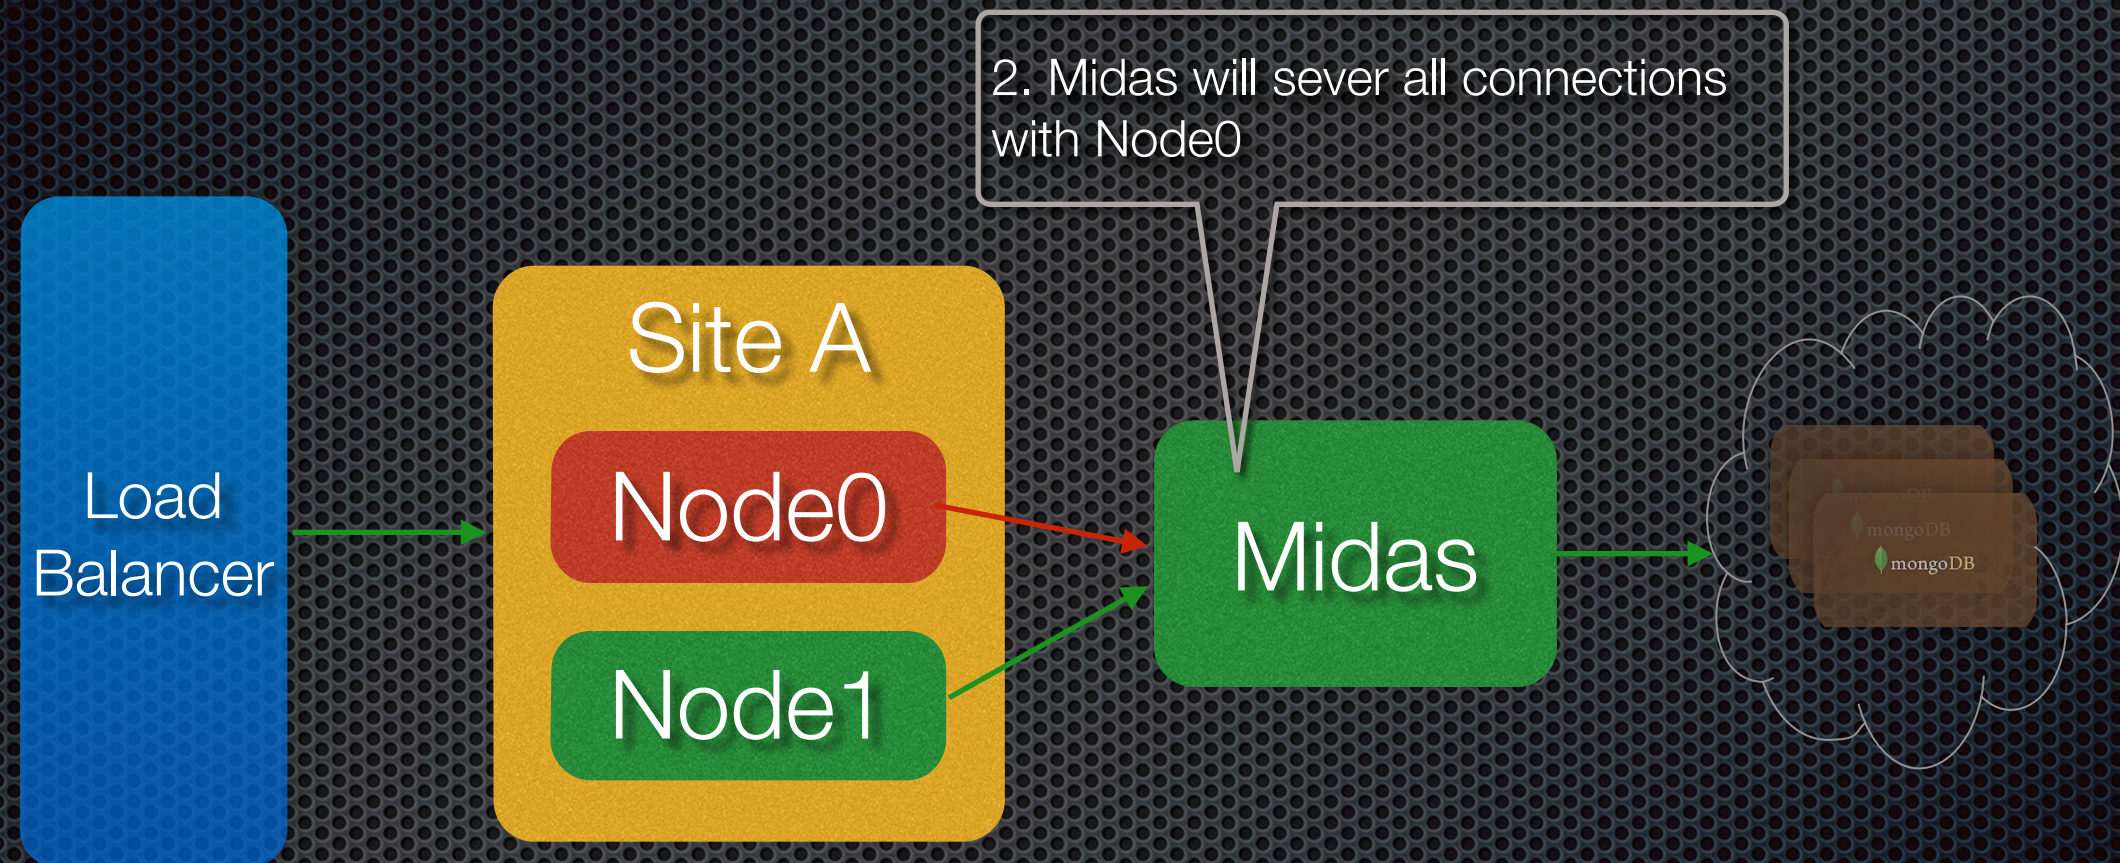
2. Remove Node0 from **app.midas**

```
demoAppV1 {
  mode = contraction
  // siteANode0 {
  //   ip = x.x.x.x
  //   changeSet = 3
  // }
  siteANode1 {
    ip = y.y.y.y
    changeSet = 3
  }
}
```



Online

Offline



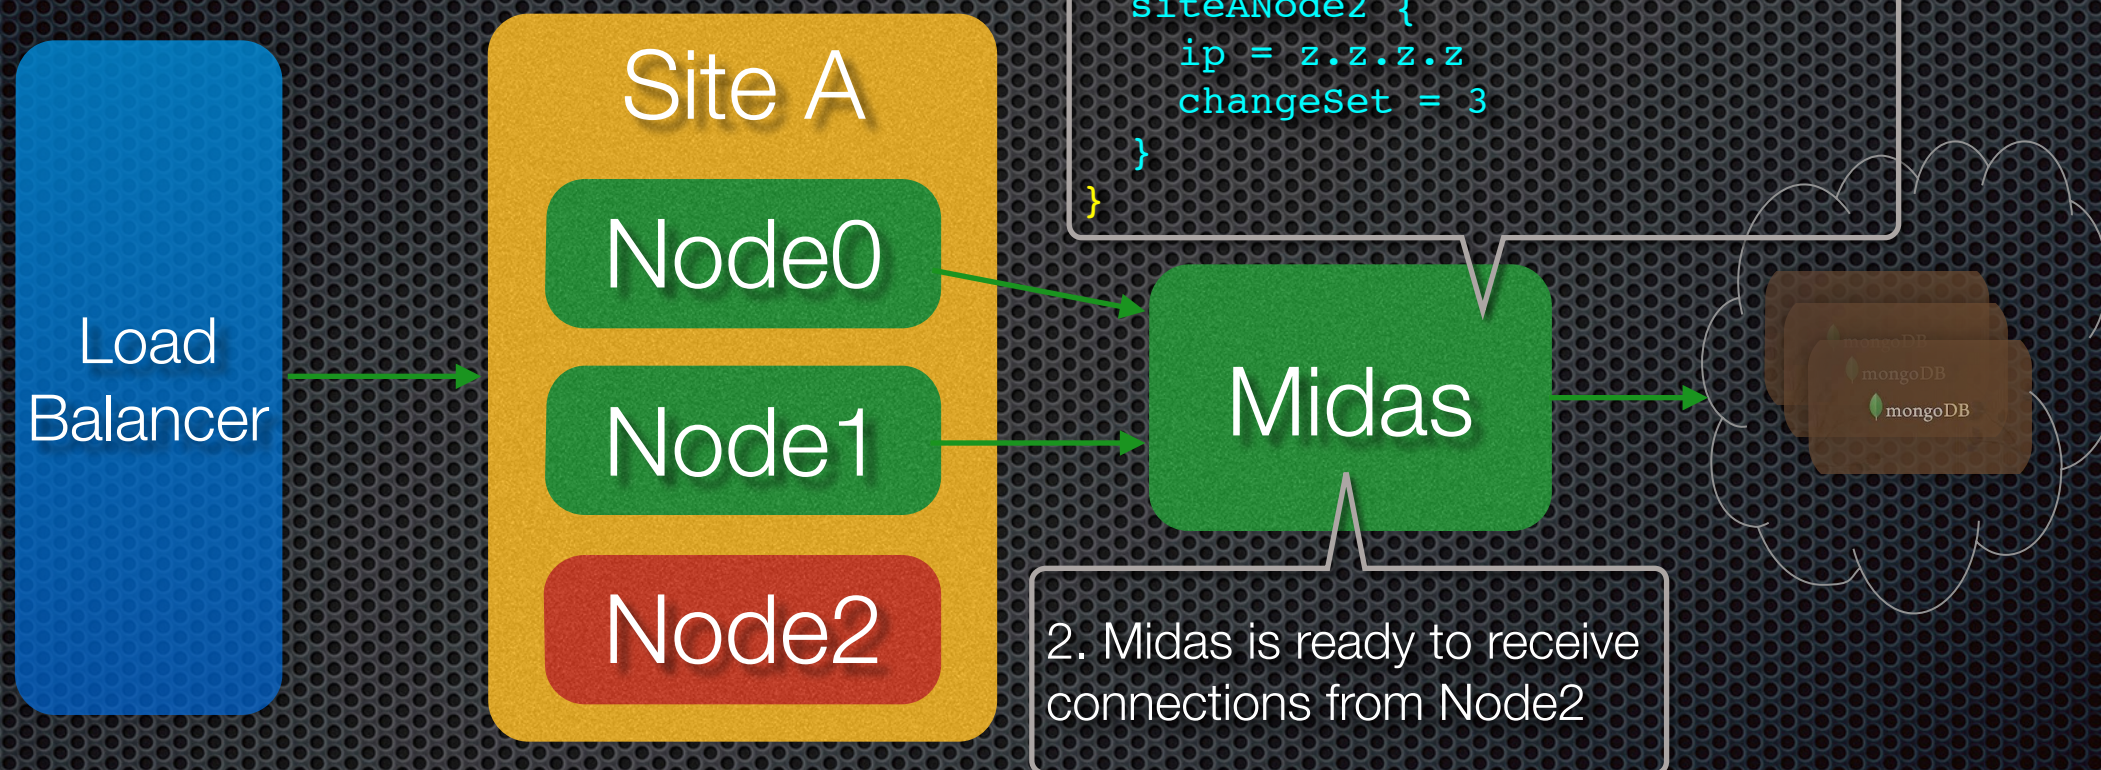
Offline does not mean node is physically taken down or is off

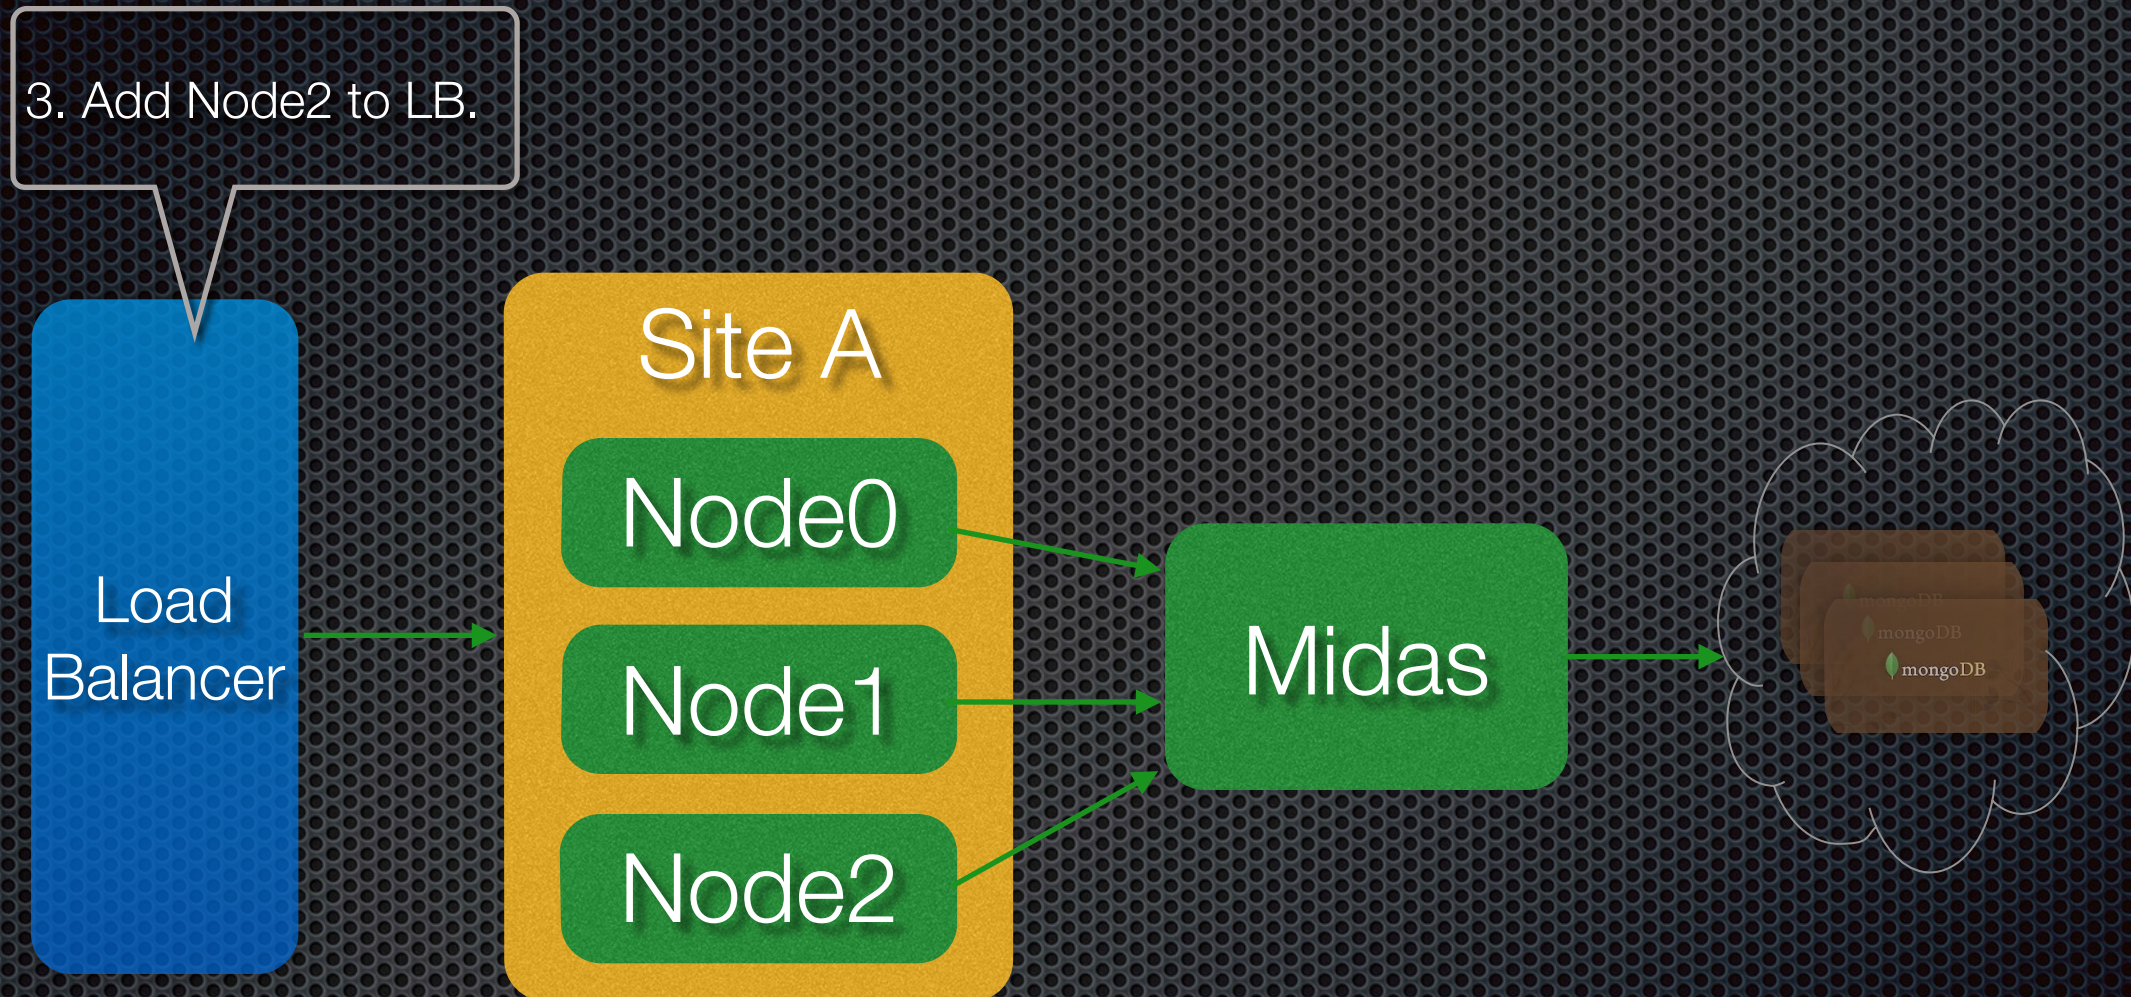
Online

Offline

Zero-Downtime Deployment Configuration - 1

Node Injection at Runtime

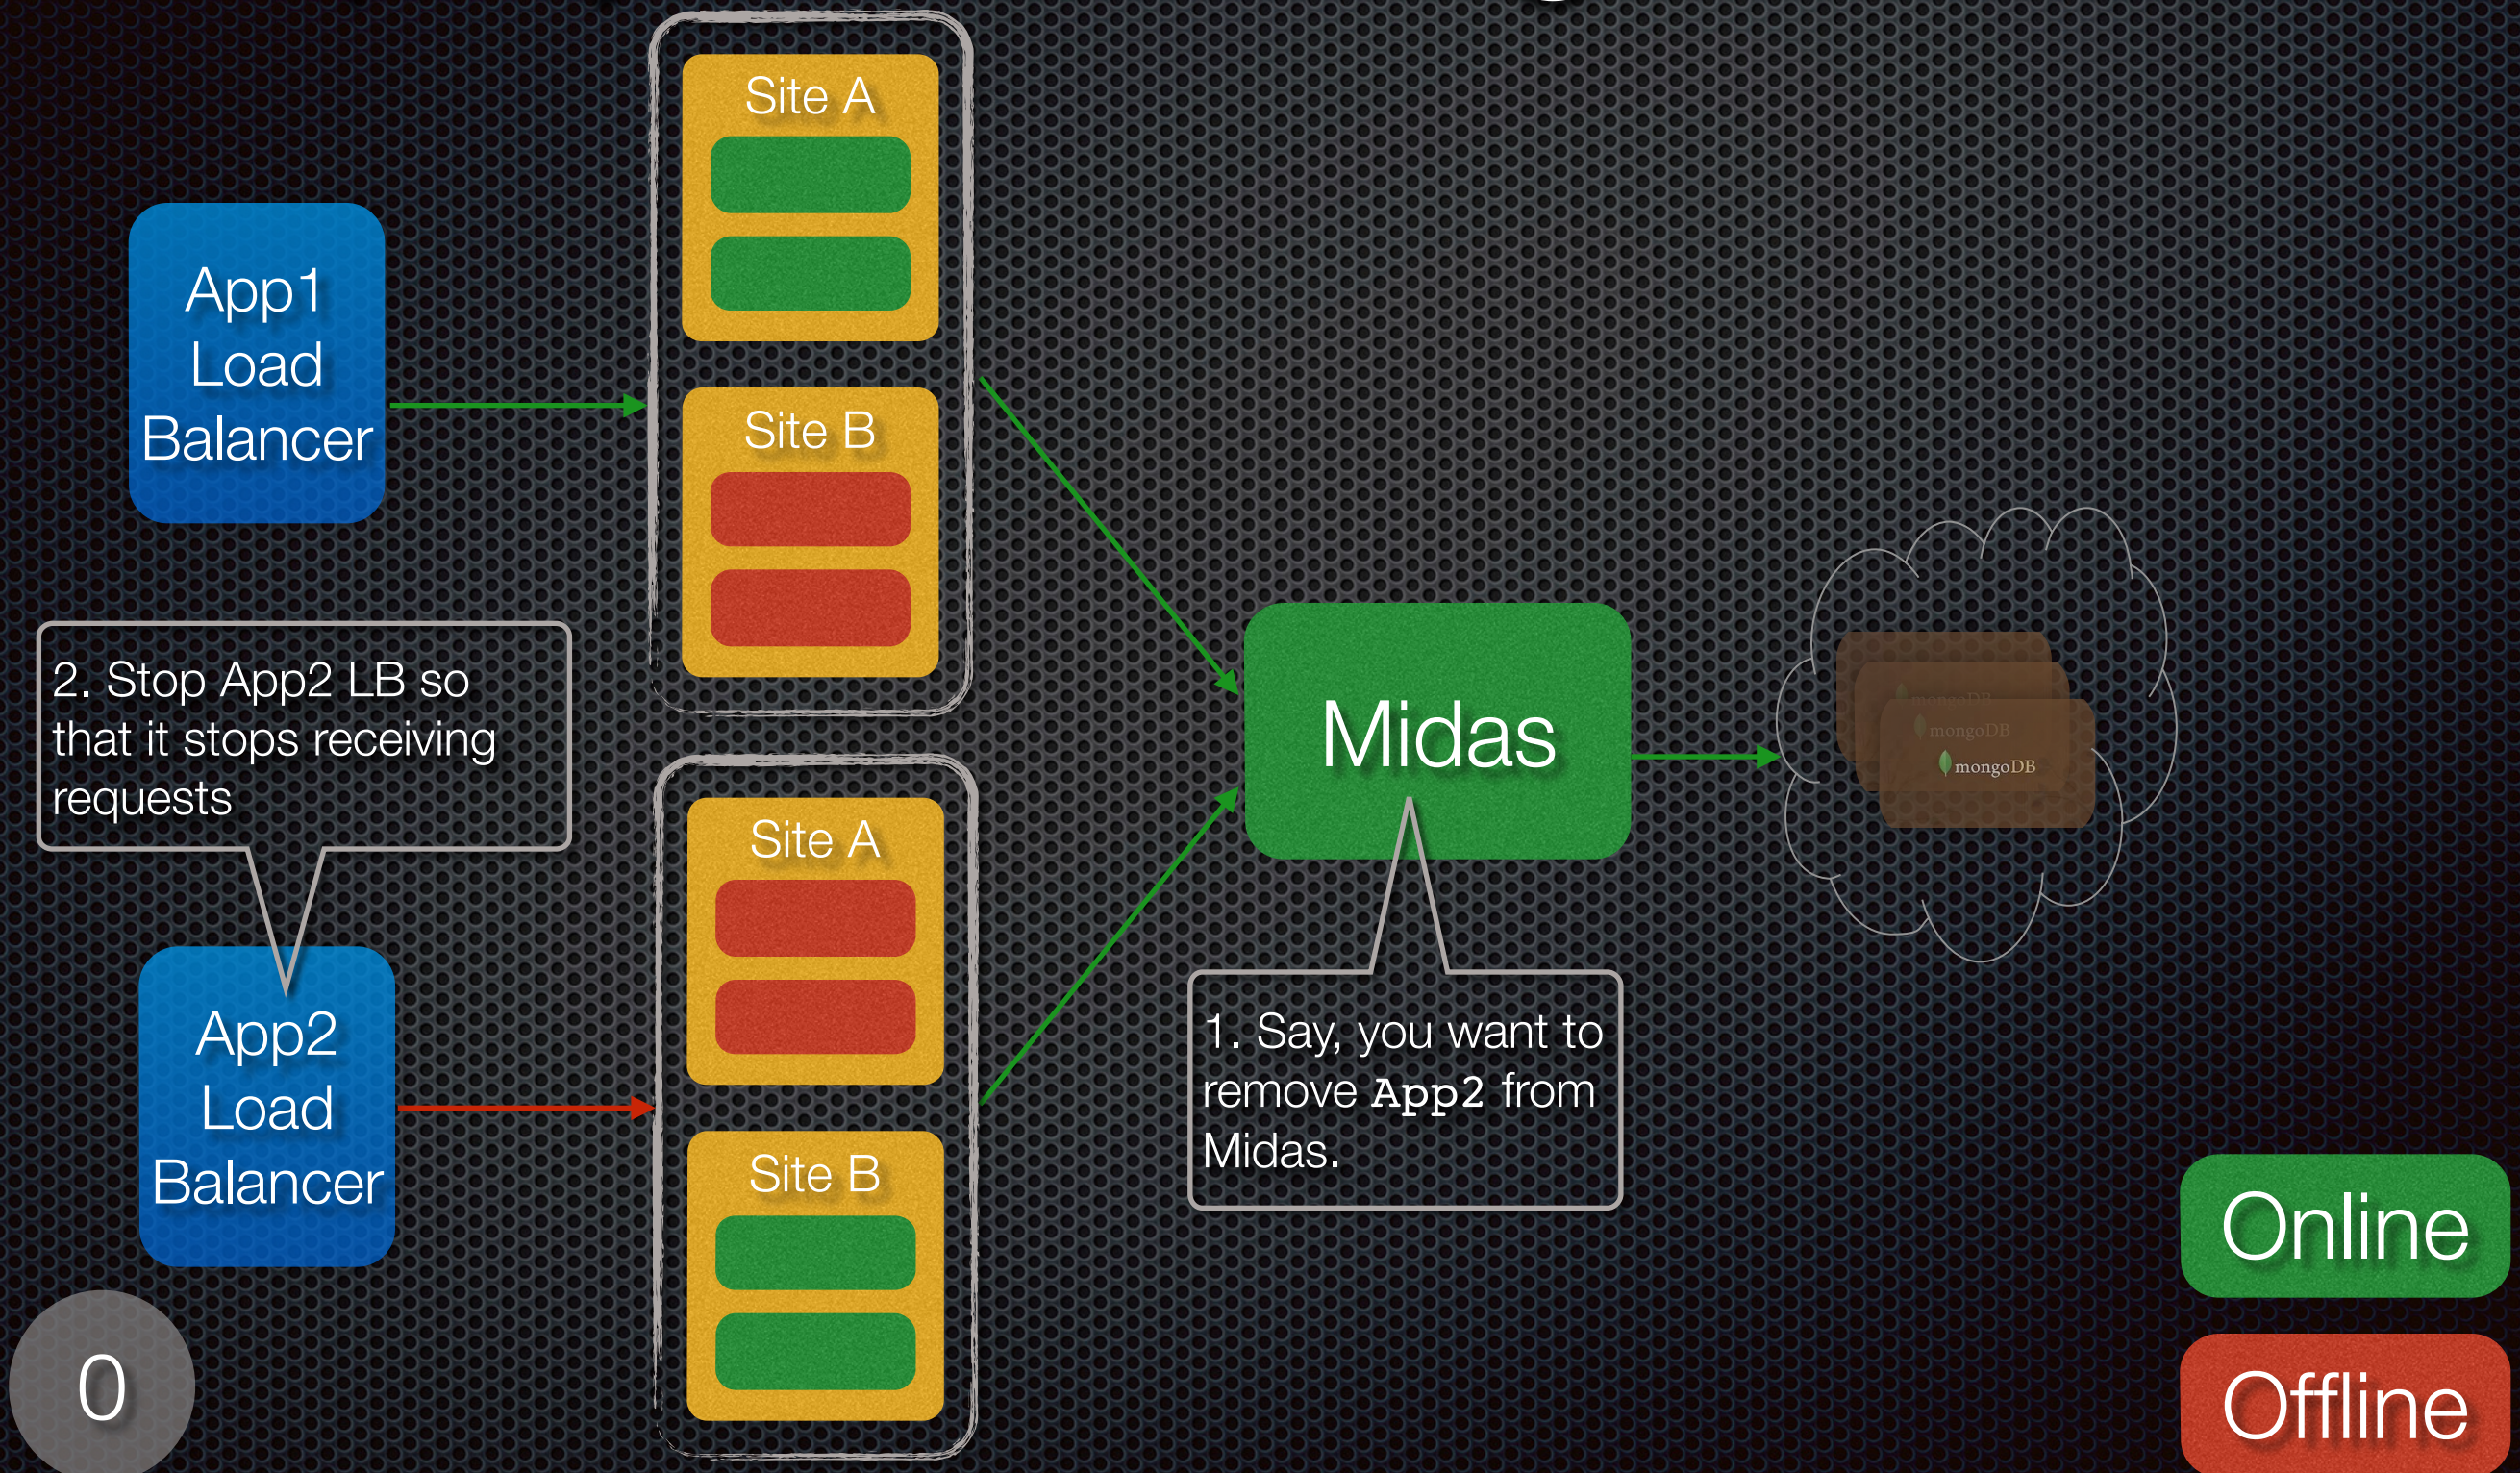


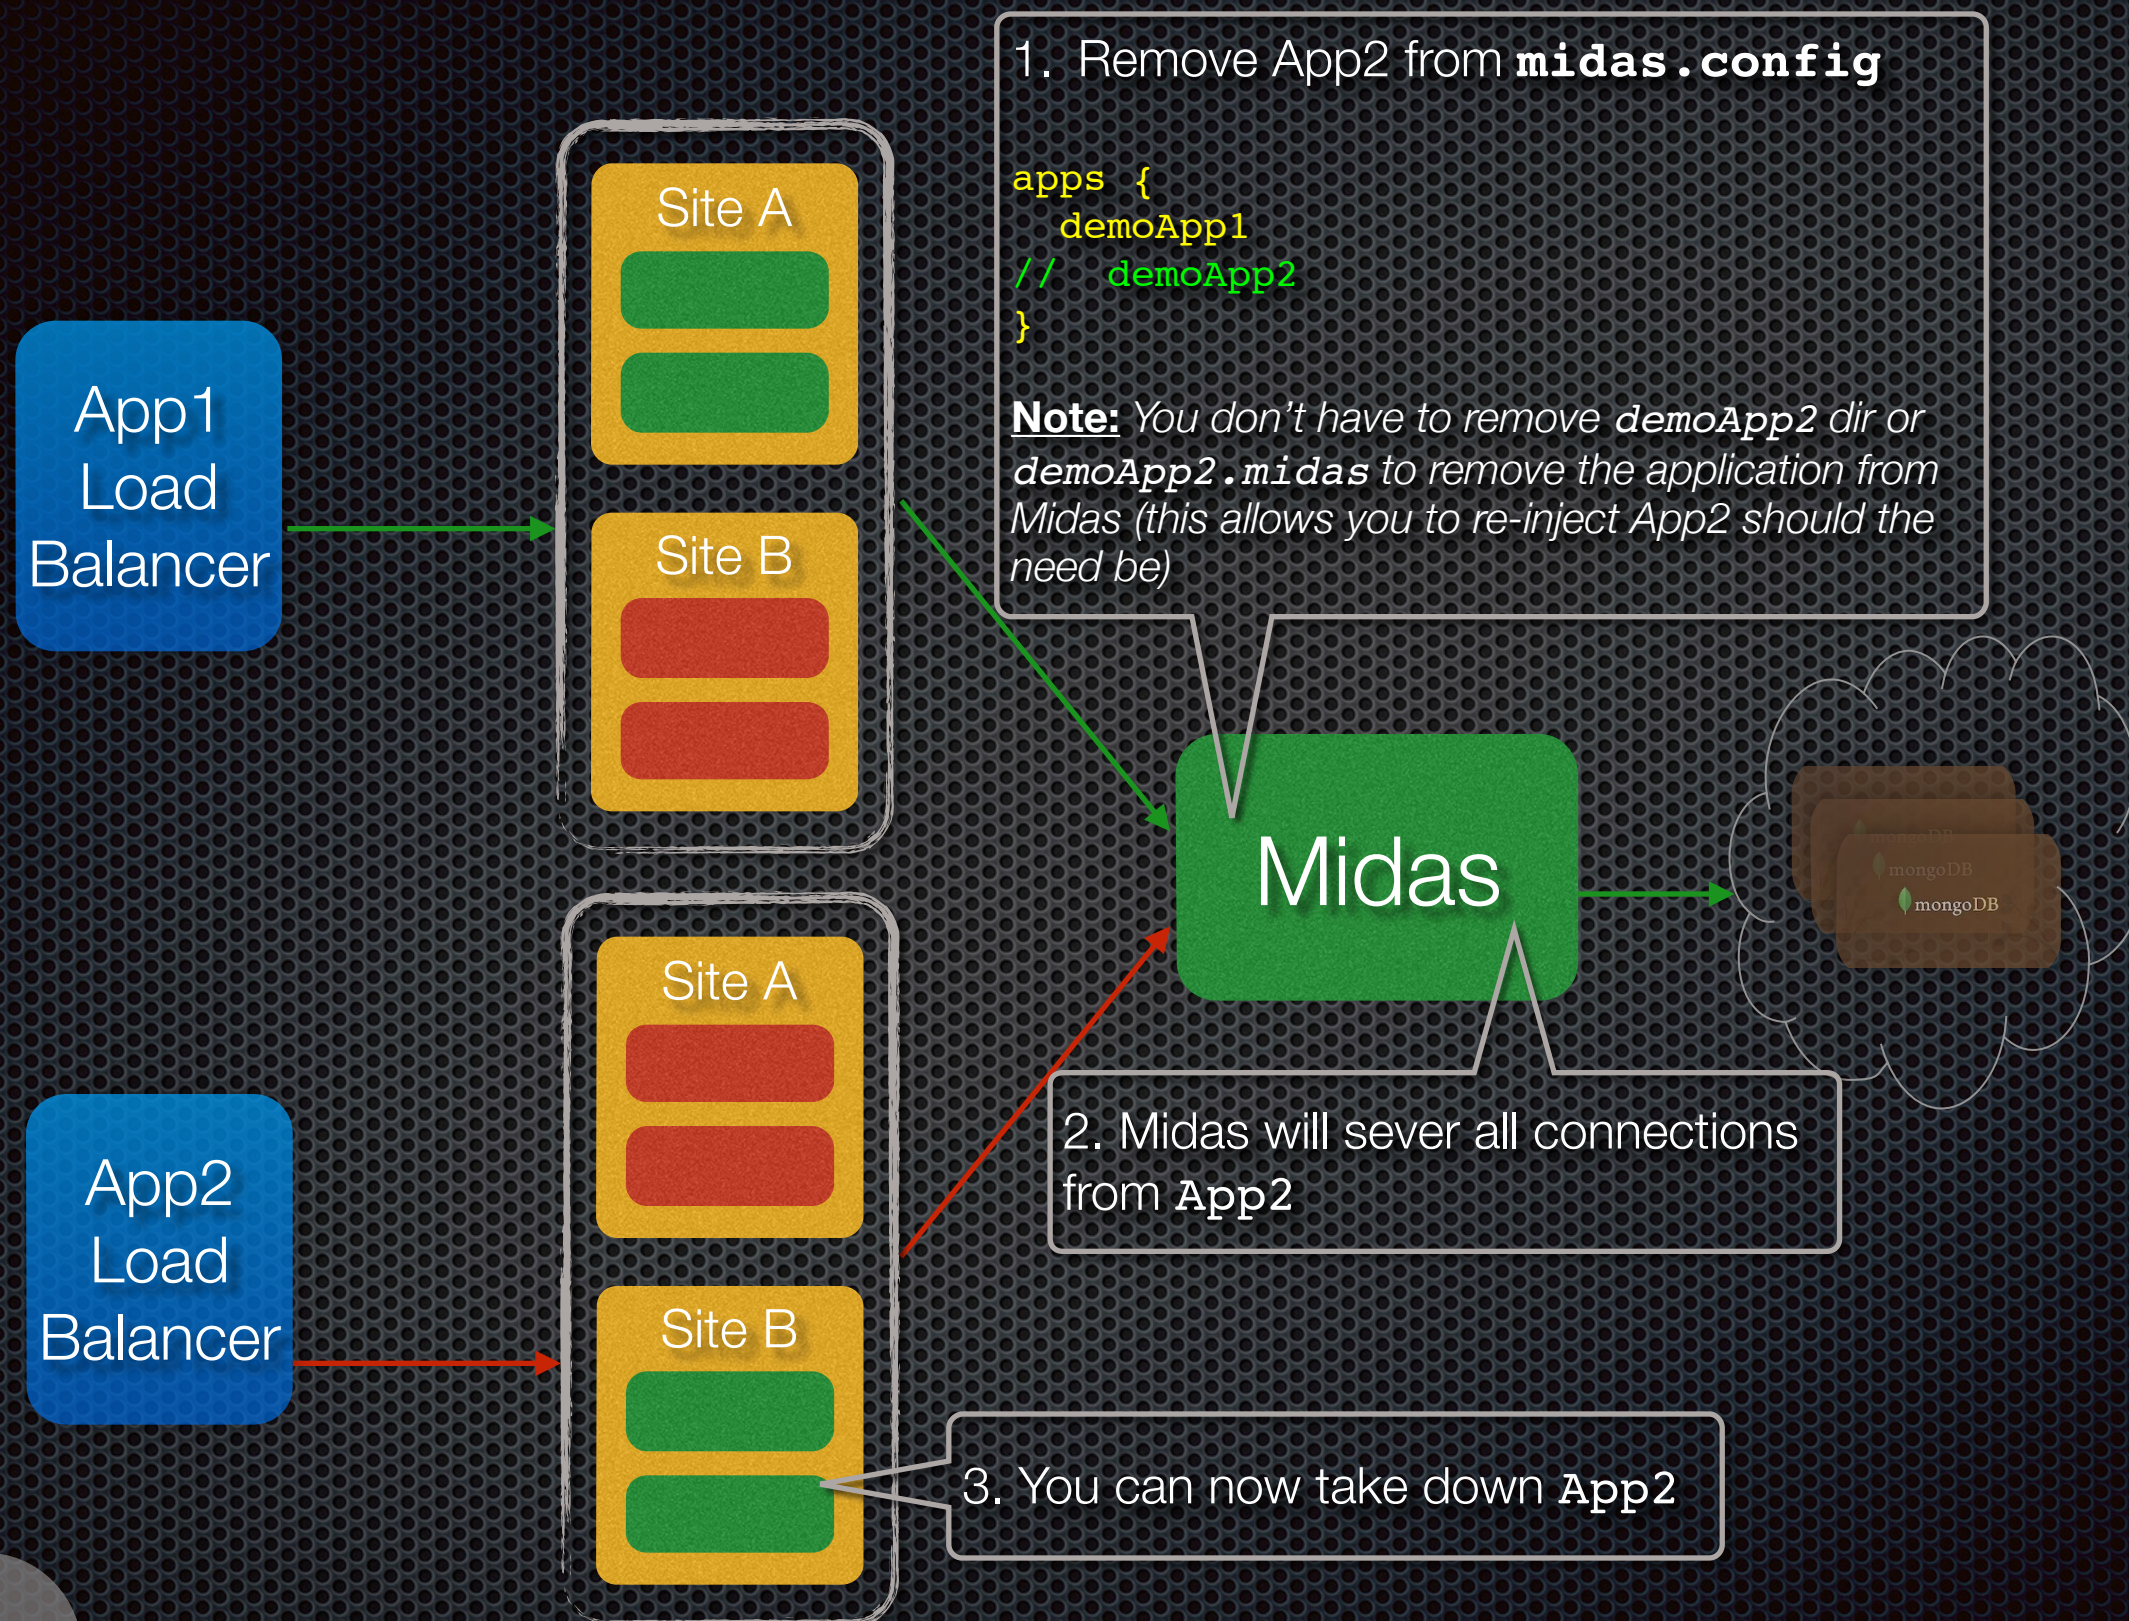


Zero-Downtime Deployment Configuration - 1

App Removal at Runtime

Deployment config - 1



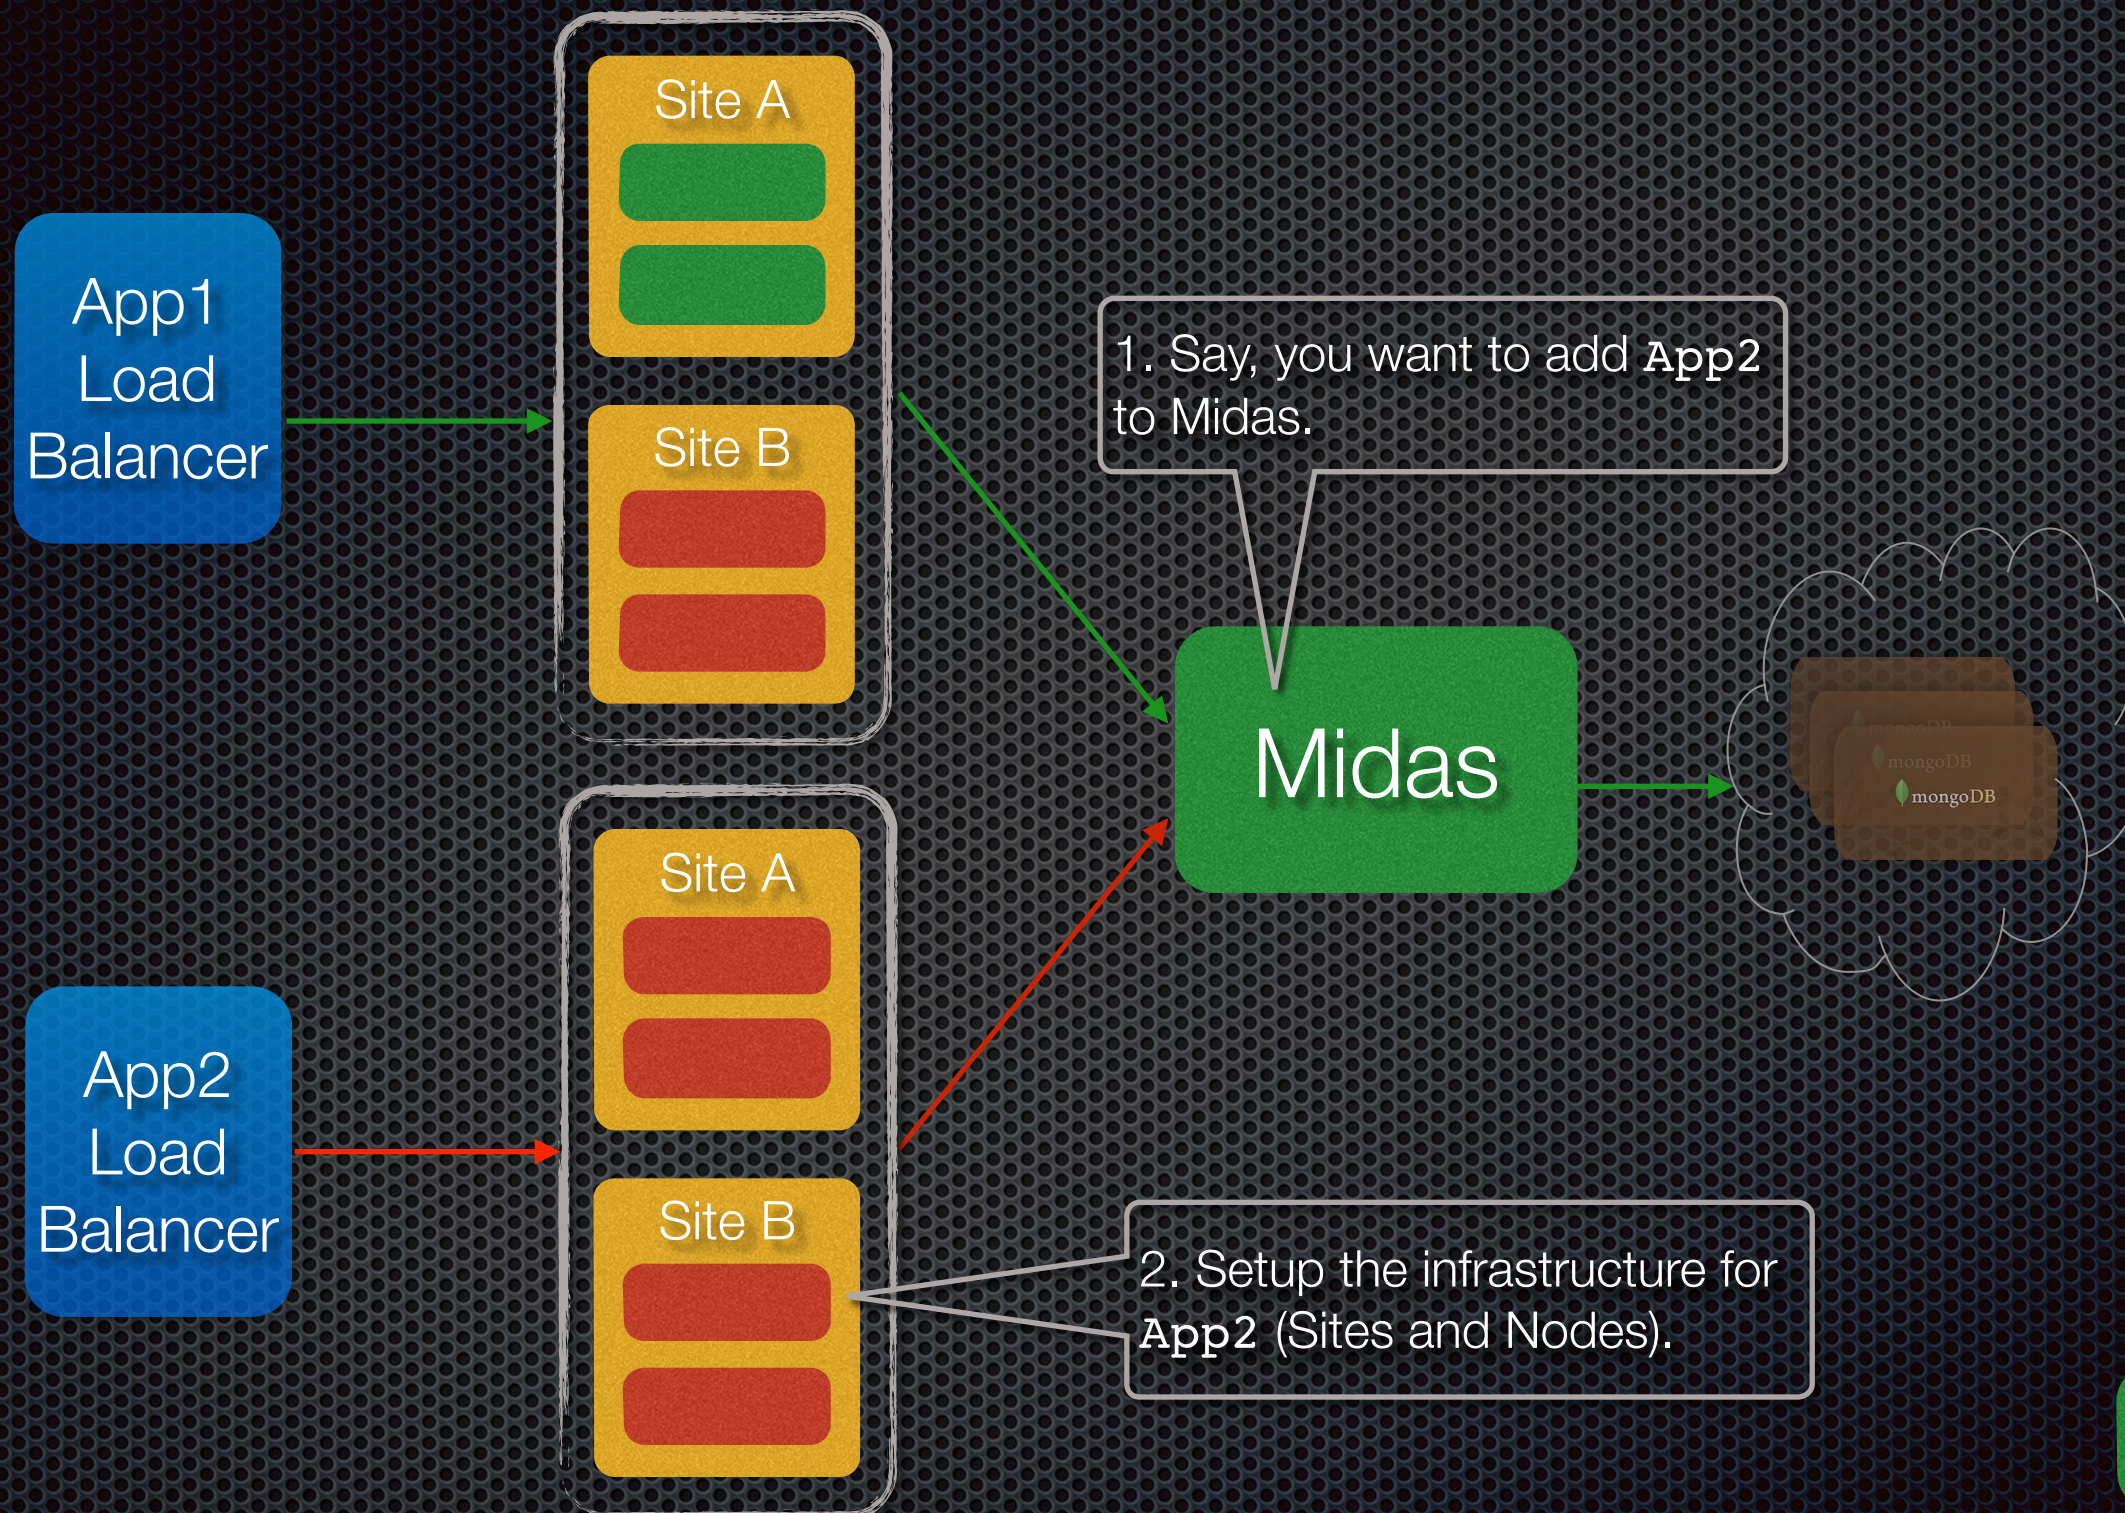


Online

Offline

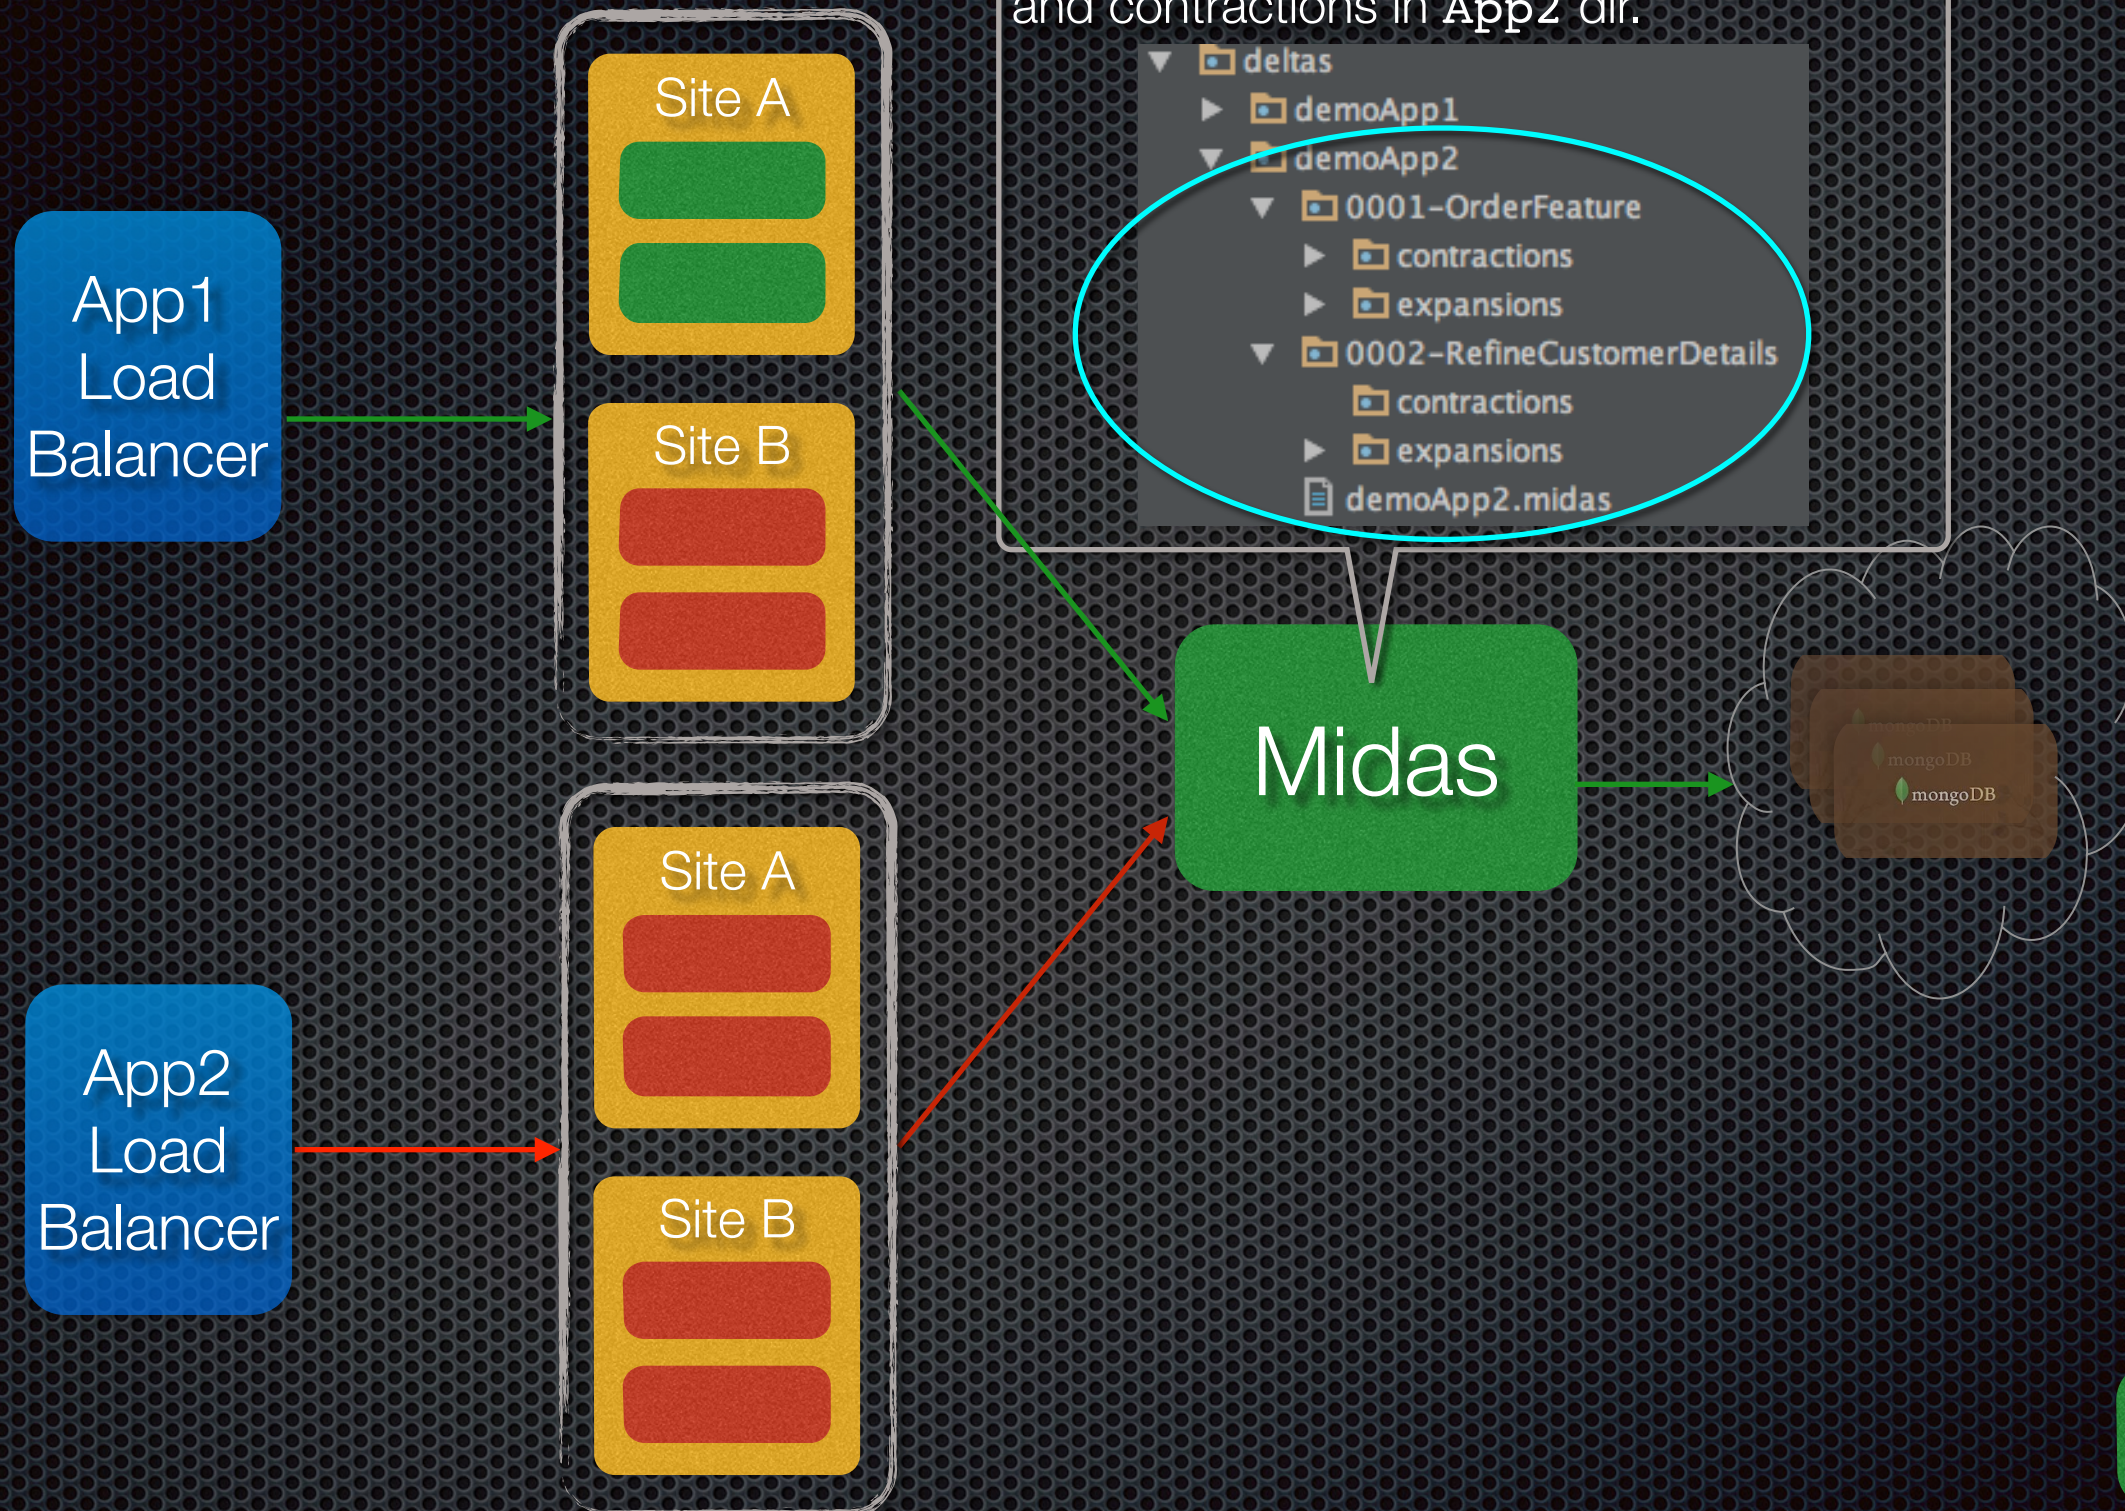
Zero-Downtime Deployment Configuration - 1

App Injection at Runtime



Online

Offline



3. Create App2 dir within deltas dir
4. Set-up change sets with expansions and contractions in App2 dir.

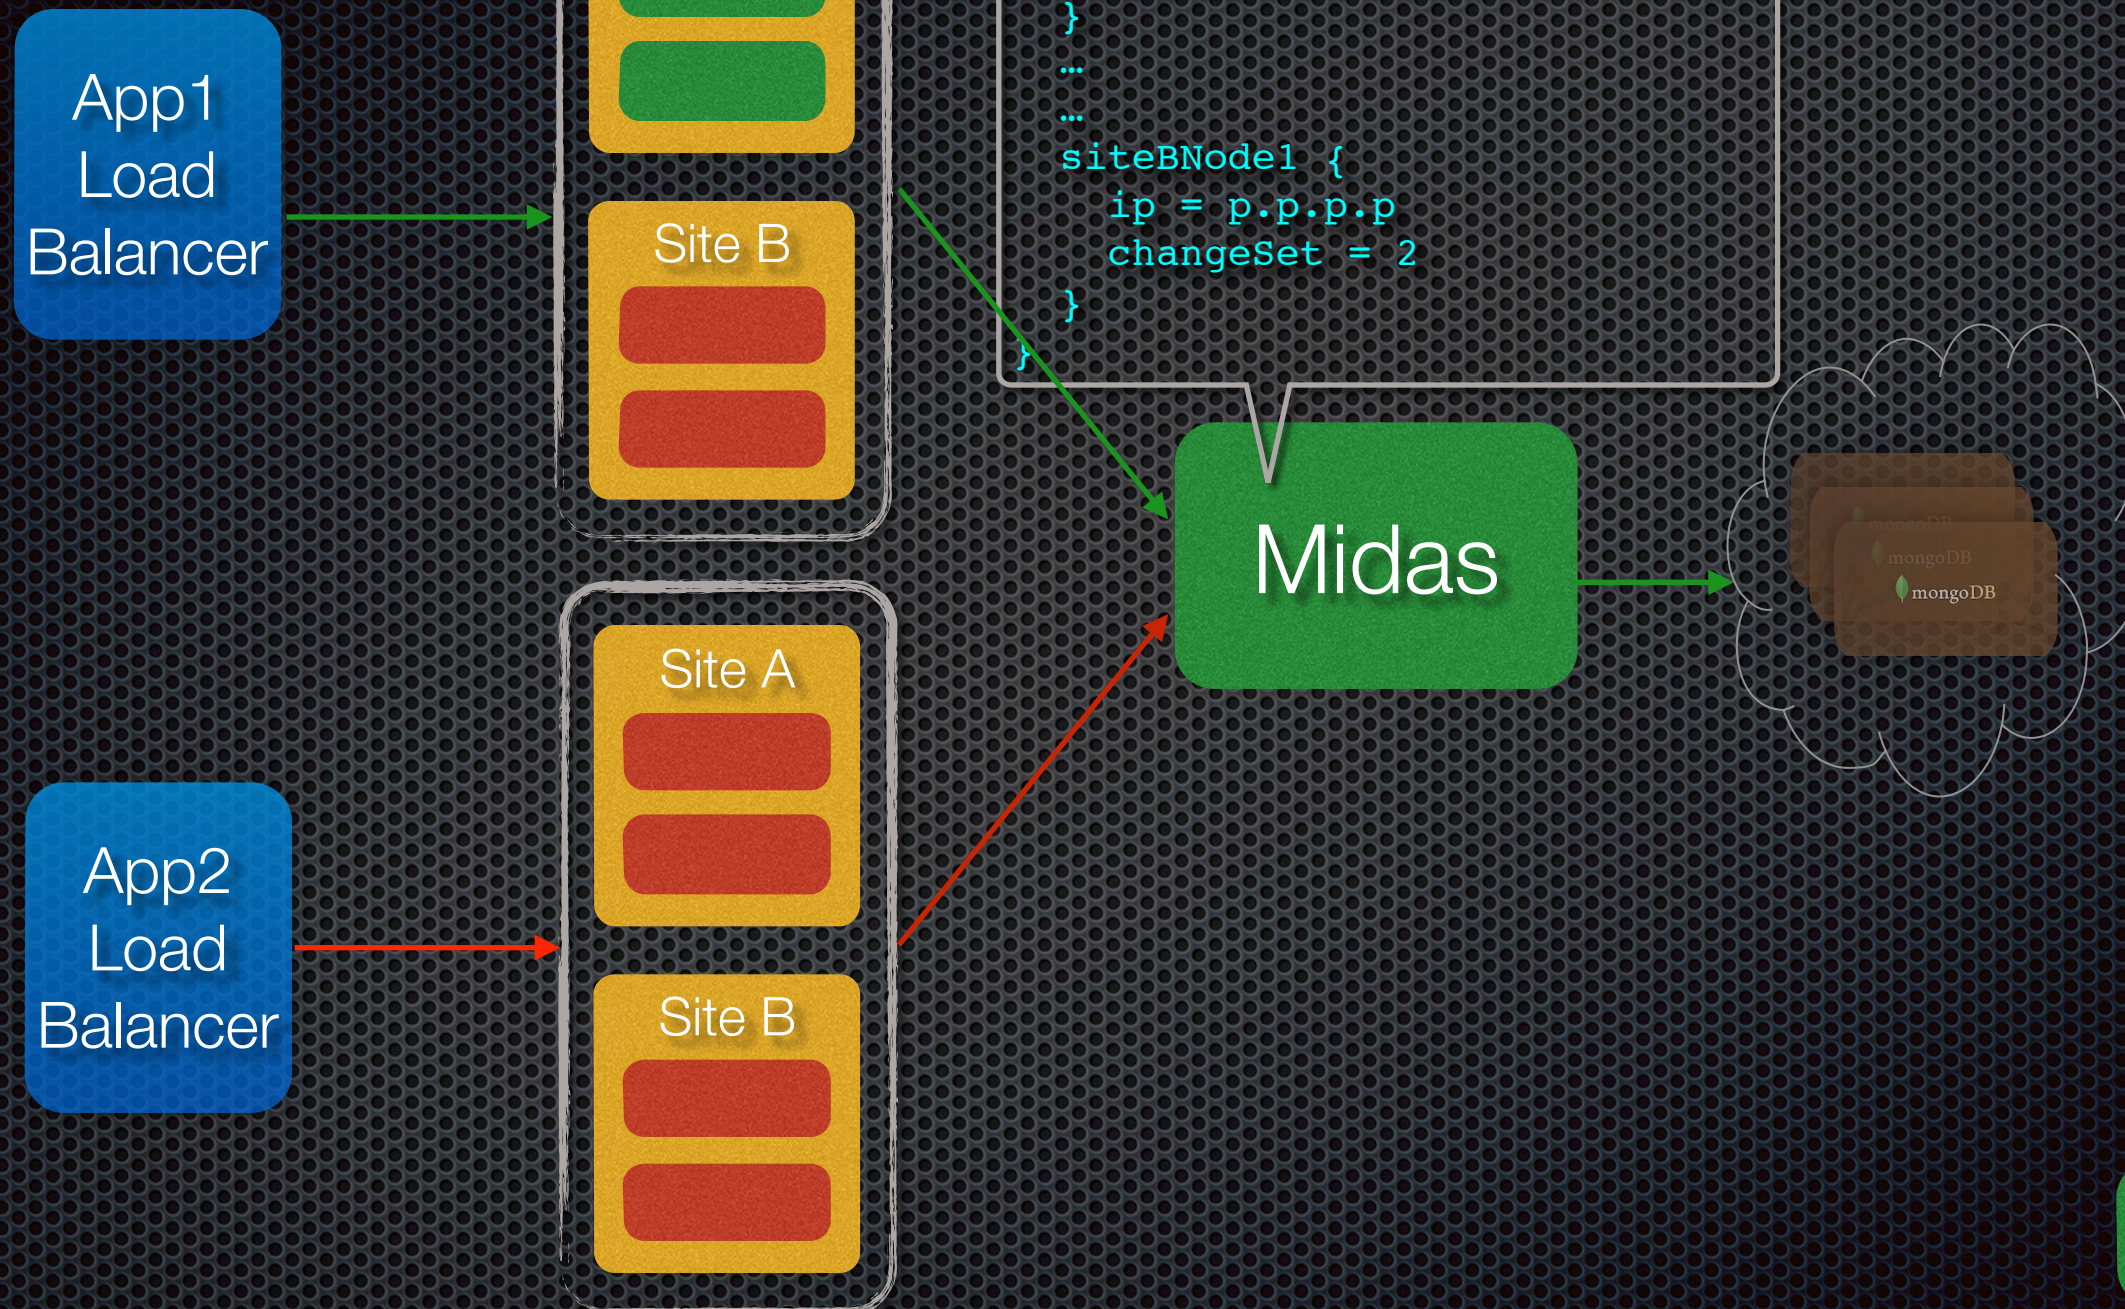
```

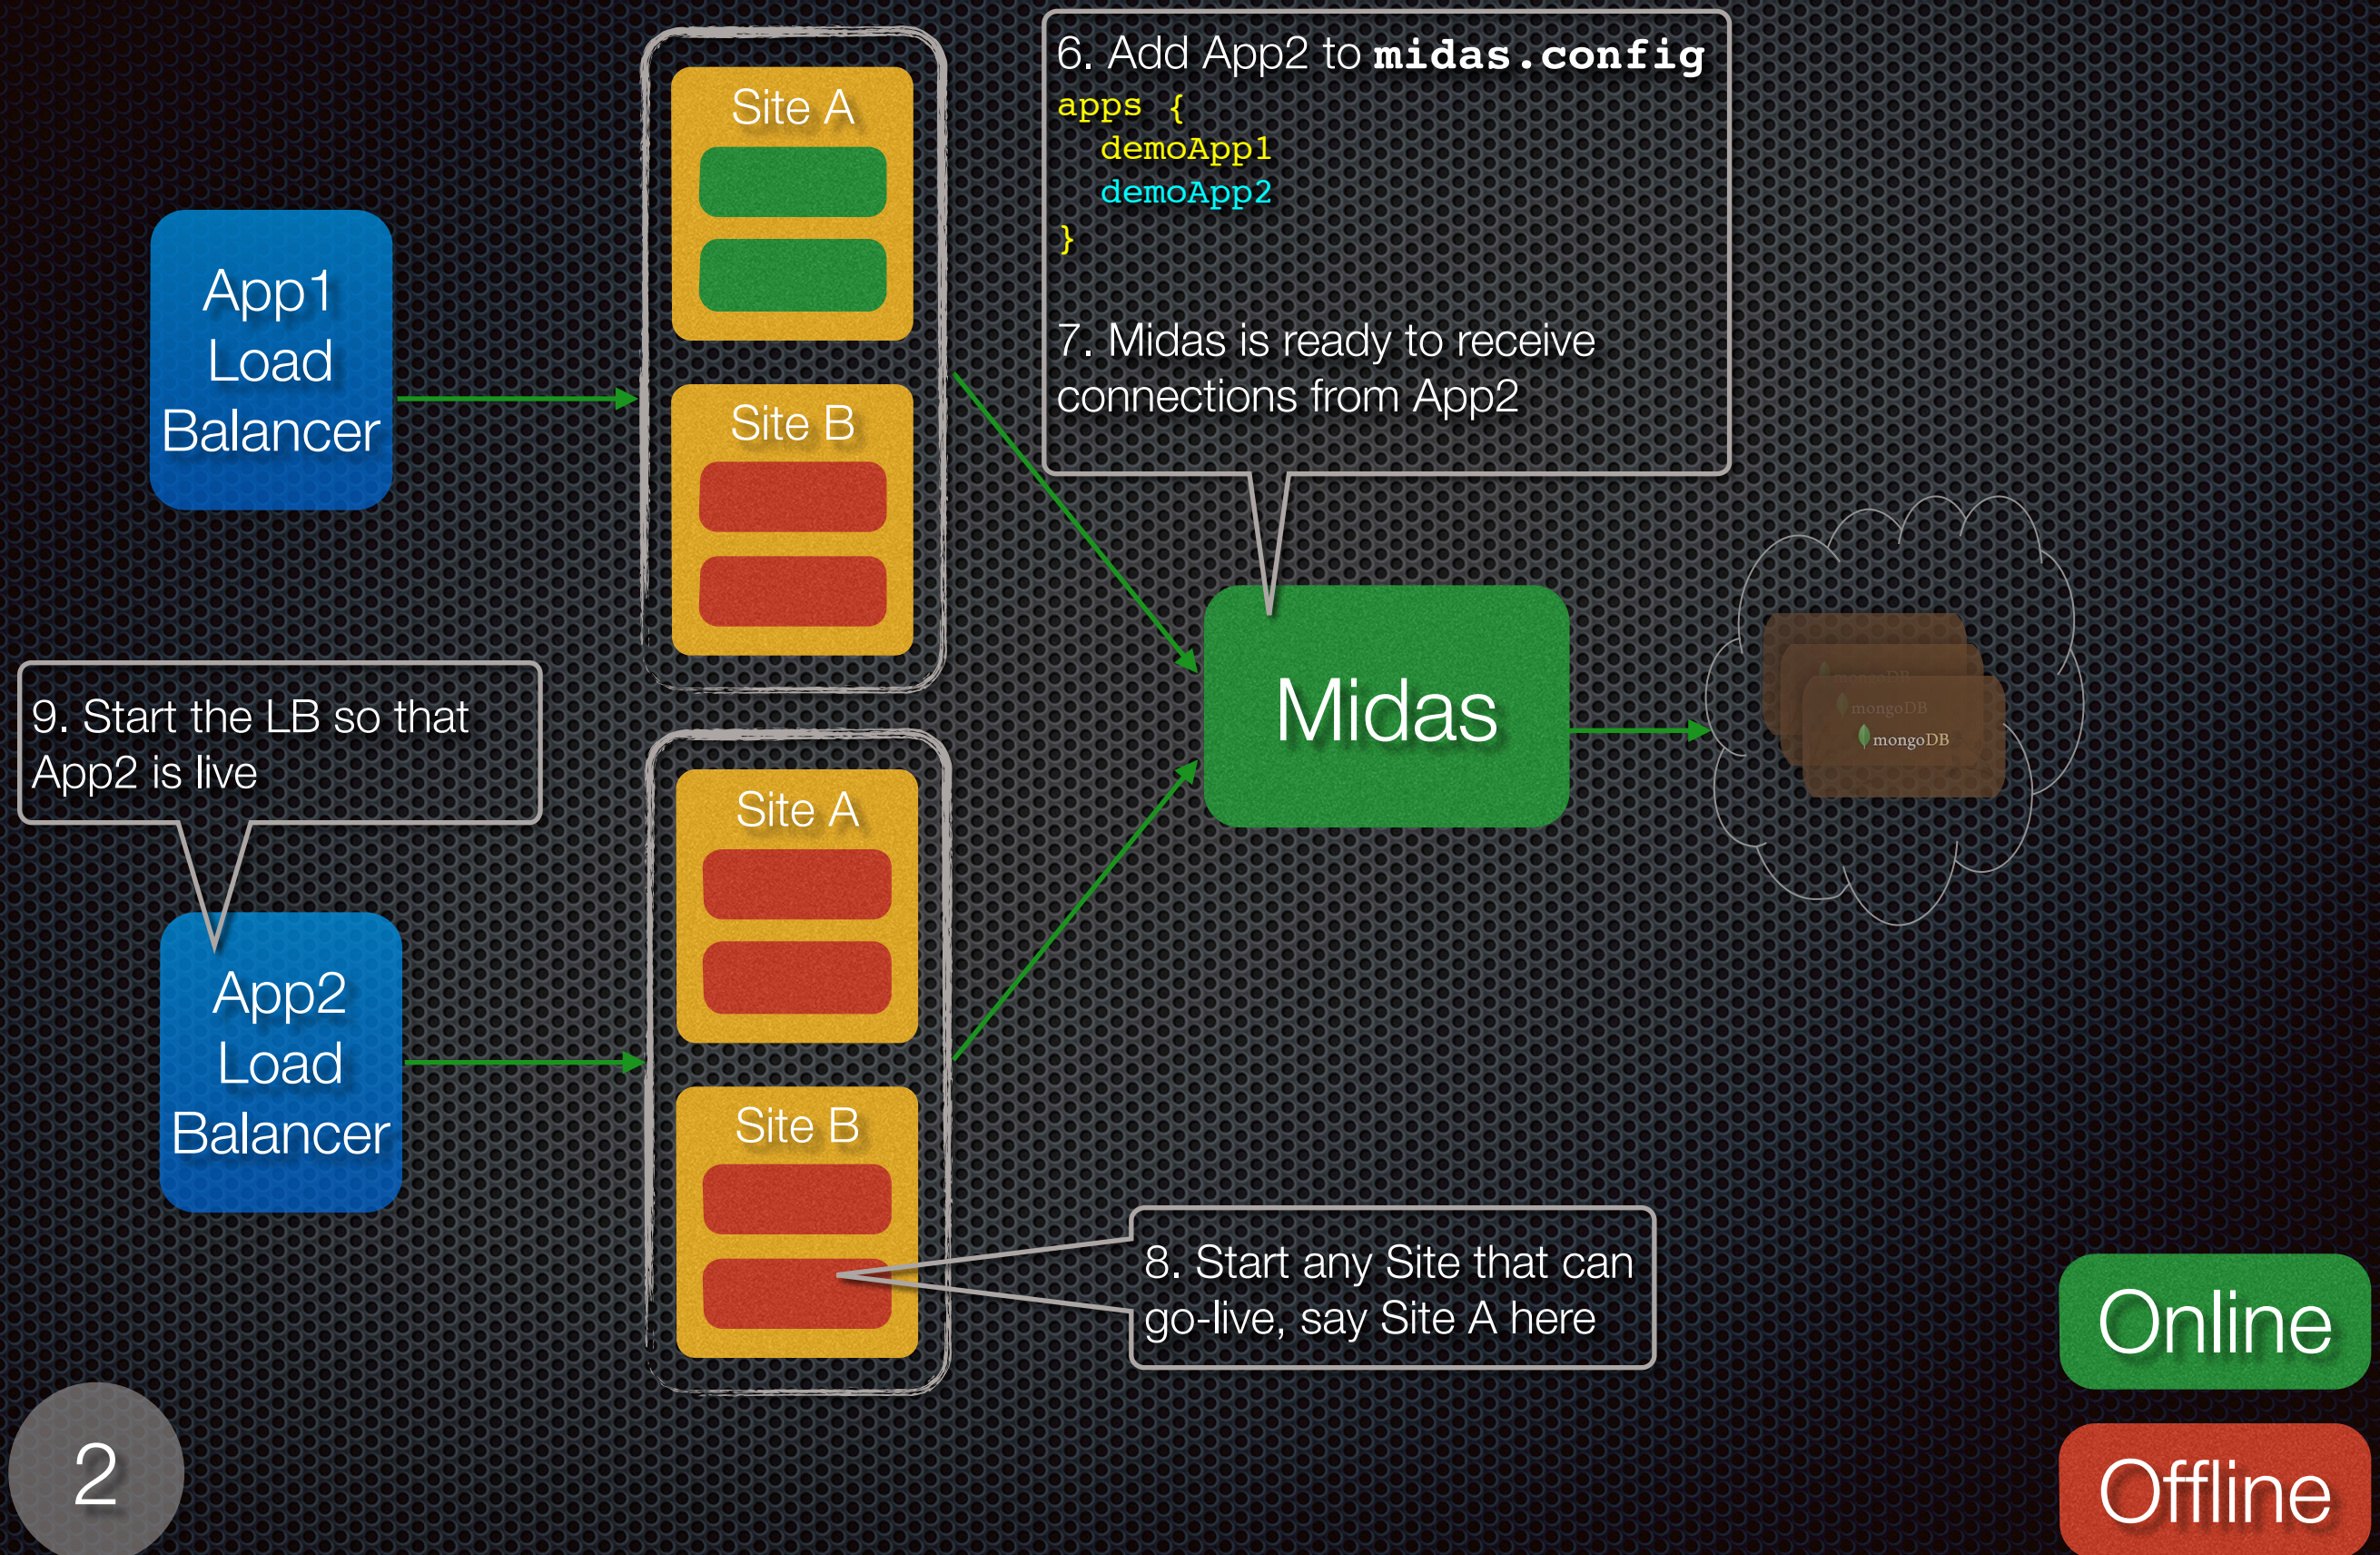
    deltas
    ├── demoApp1
    └── demoApp2
        ├── 0001-OrderFeature
        │   ├── contractions
        │   └── expansions
        ├── 0002-RefineCustomerDetails
        │   ├── contractions
        │   └── expansions
        └── demoApp2.midas
  
```

Midas

Online

Offline

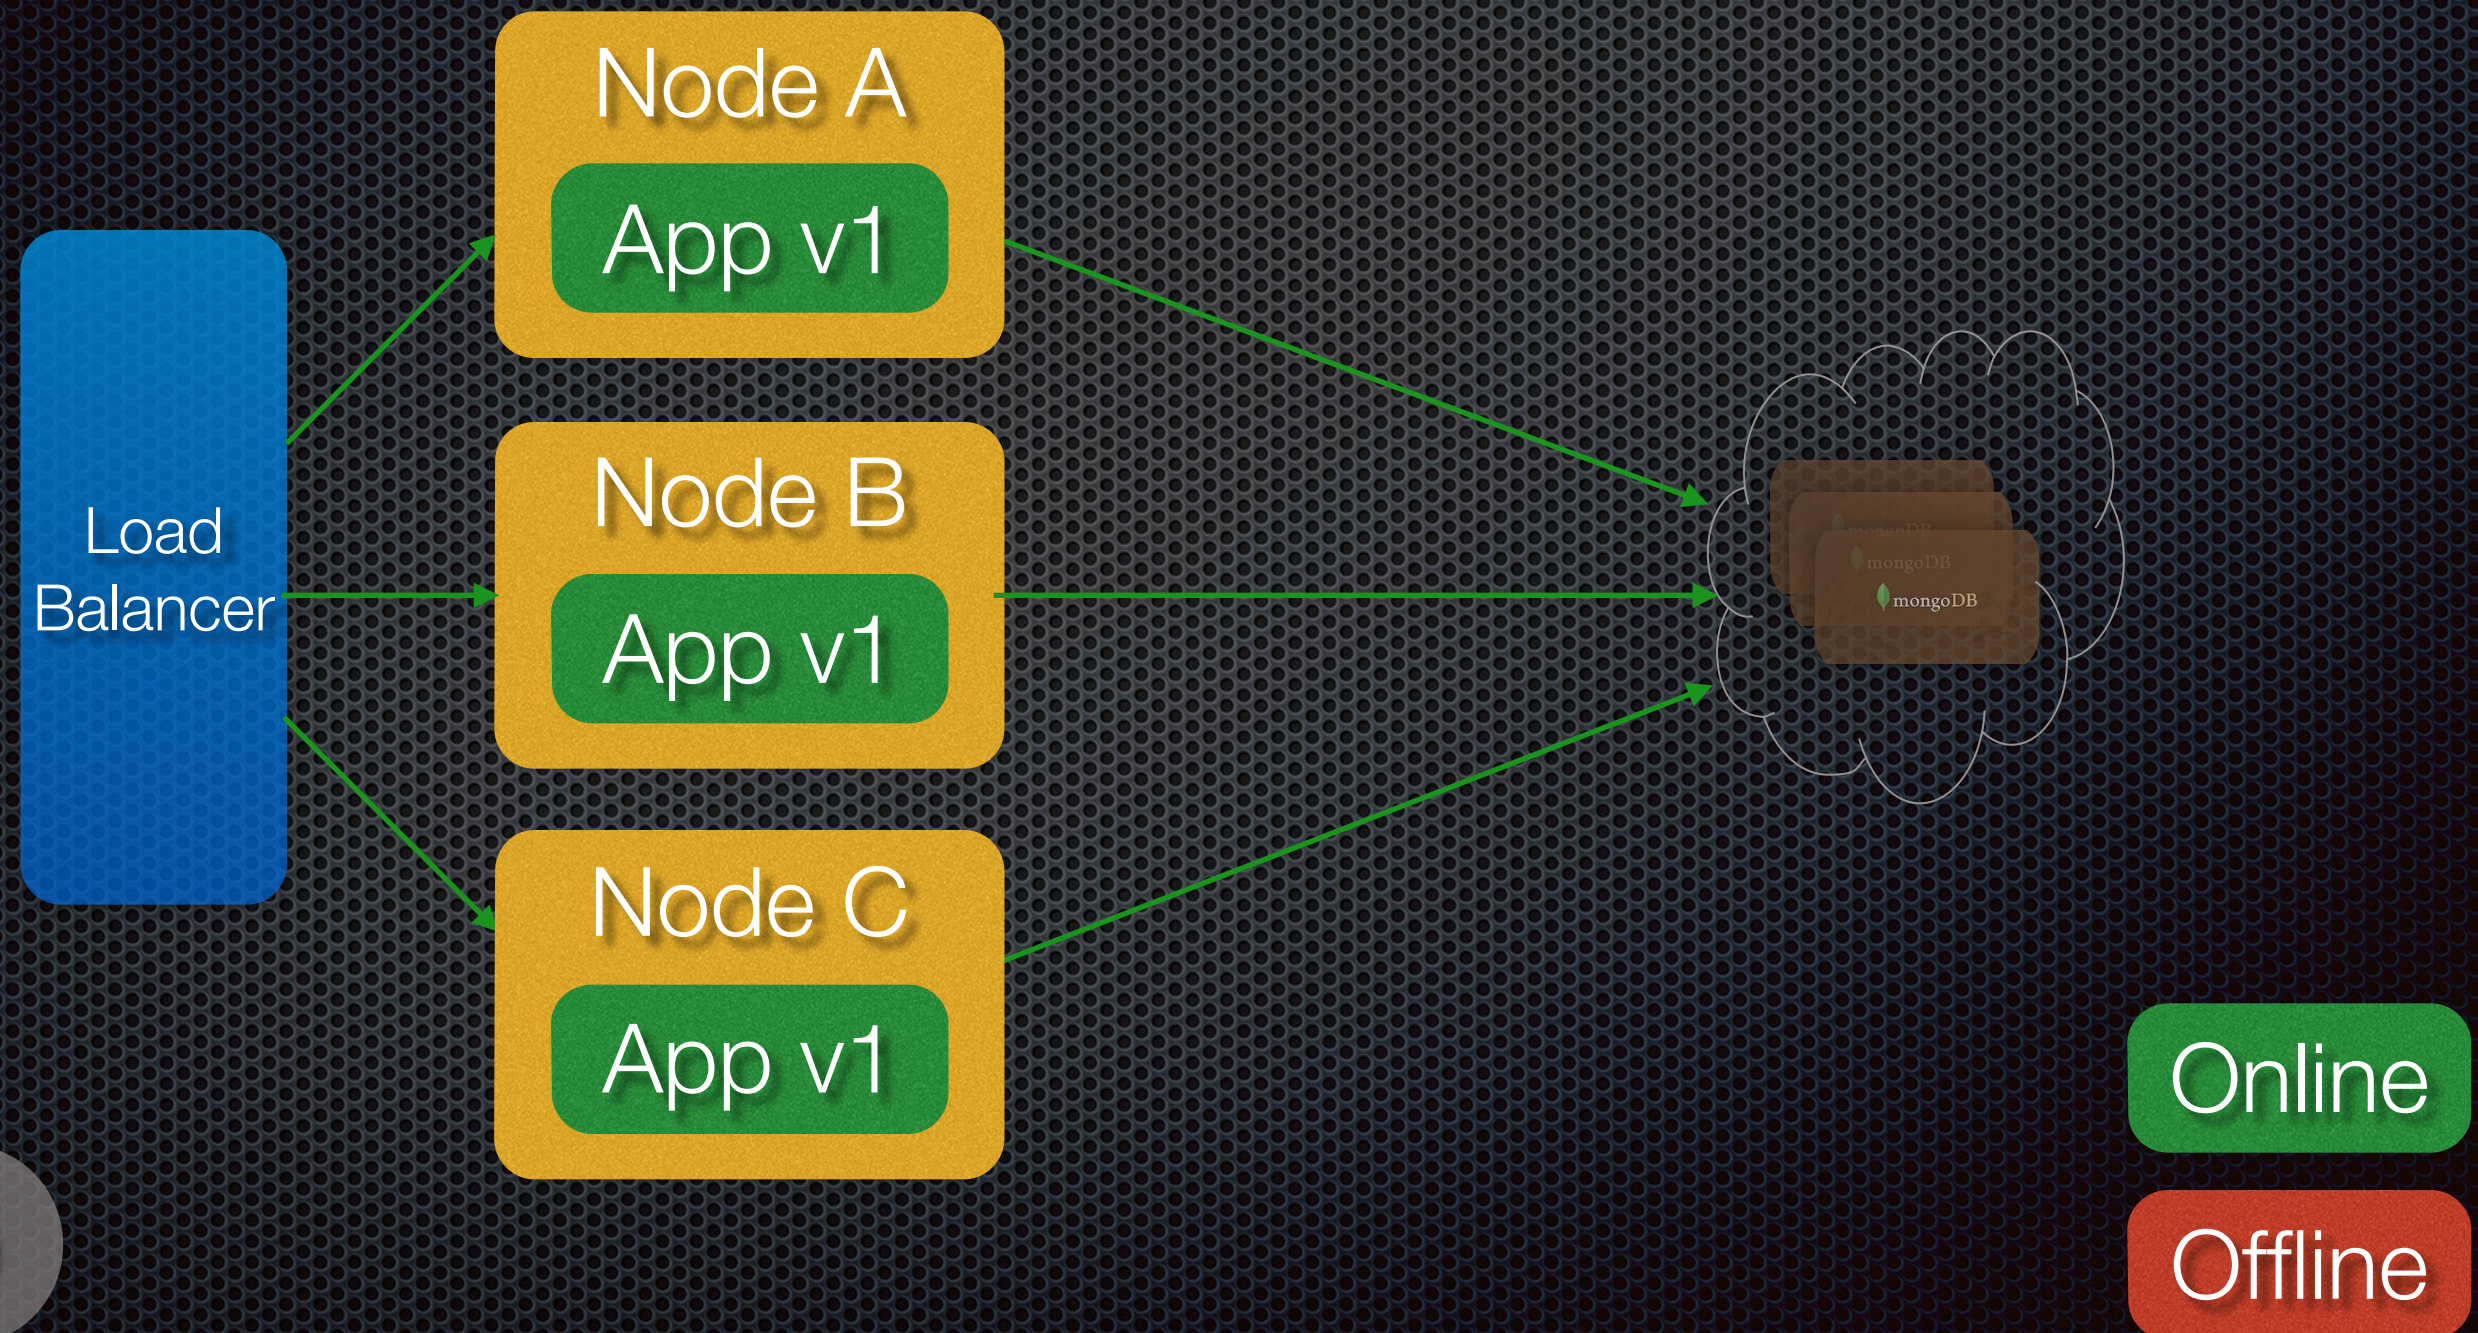




Zero-Downtime Deployment Configuration - 2

Not Supported in this Release

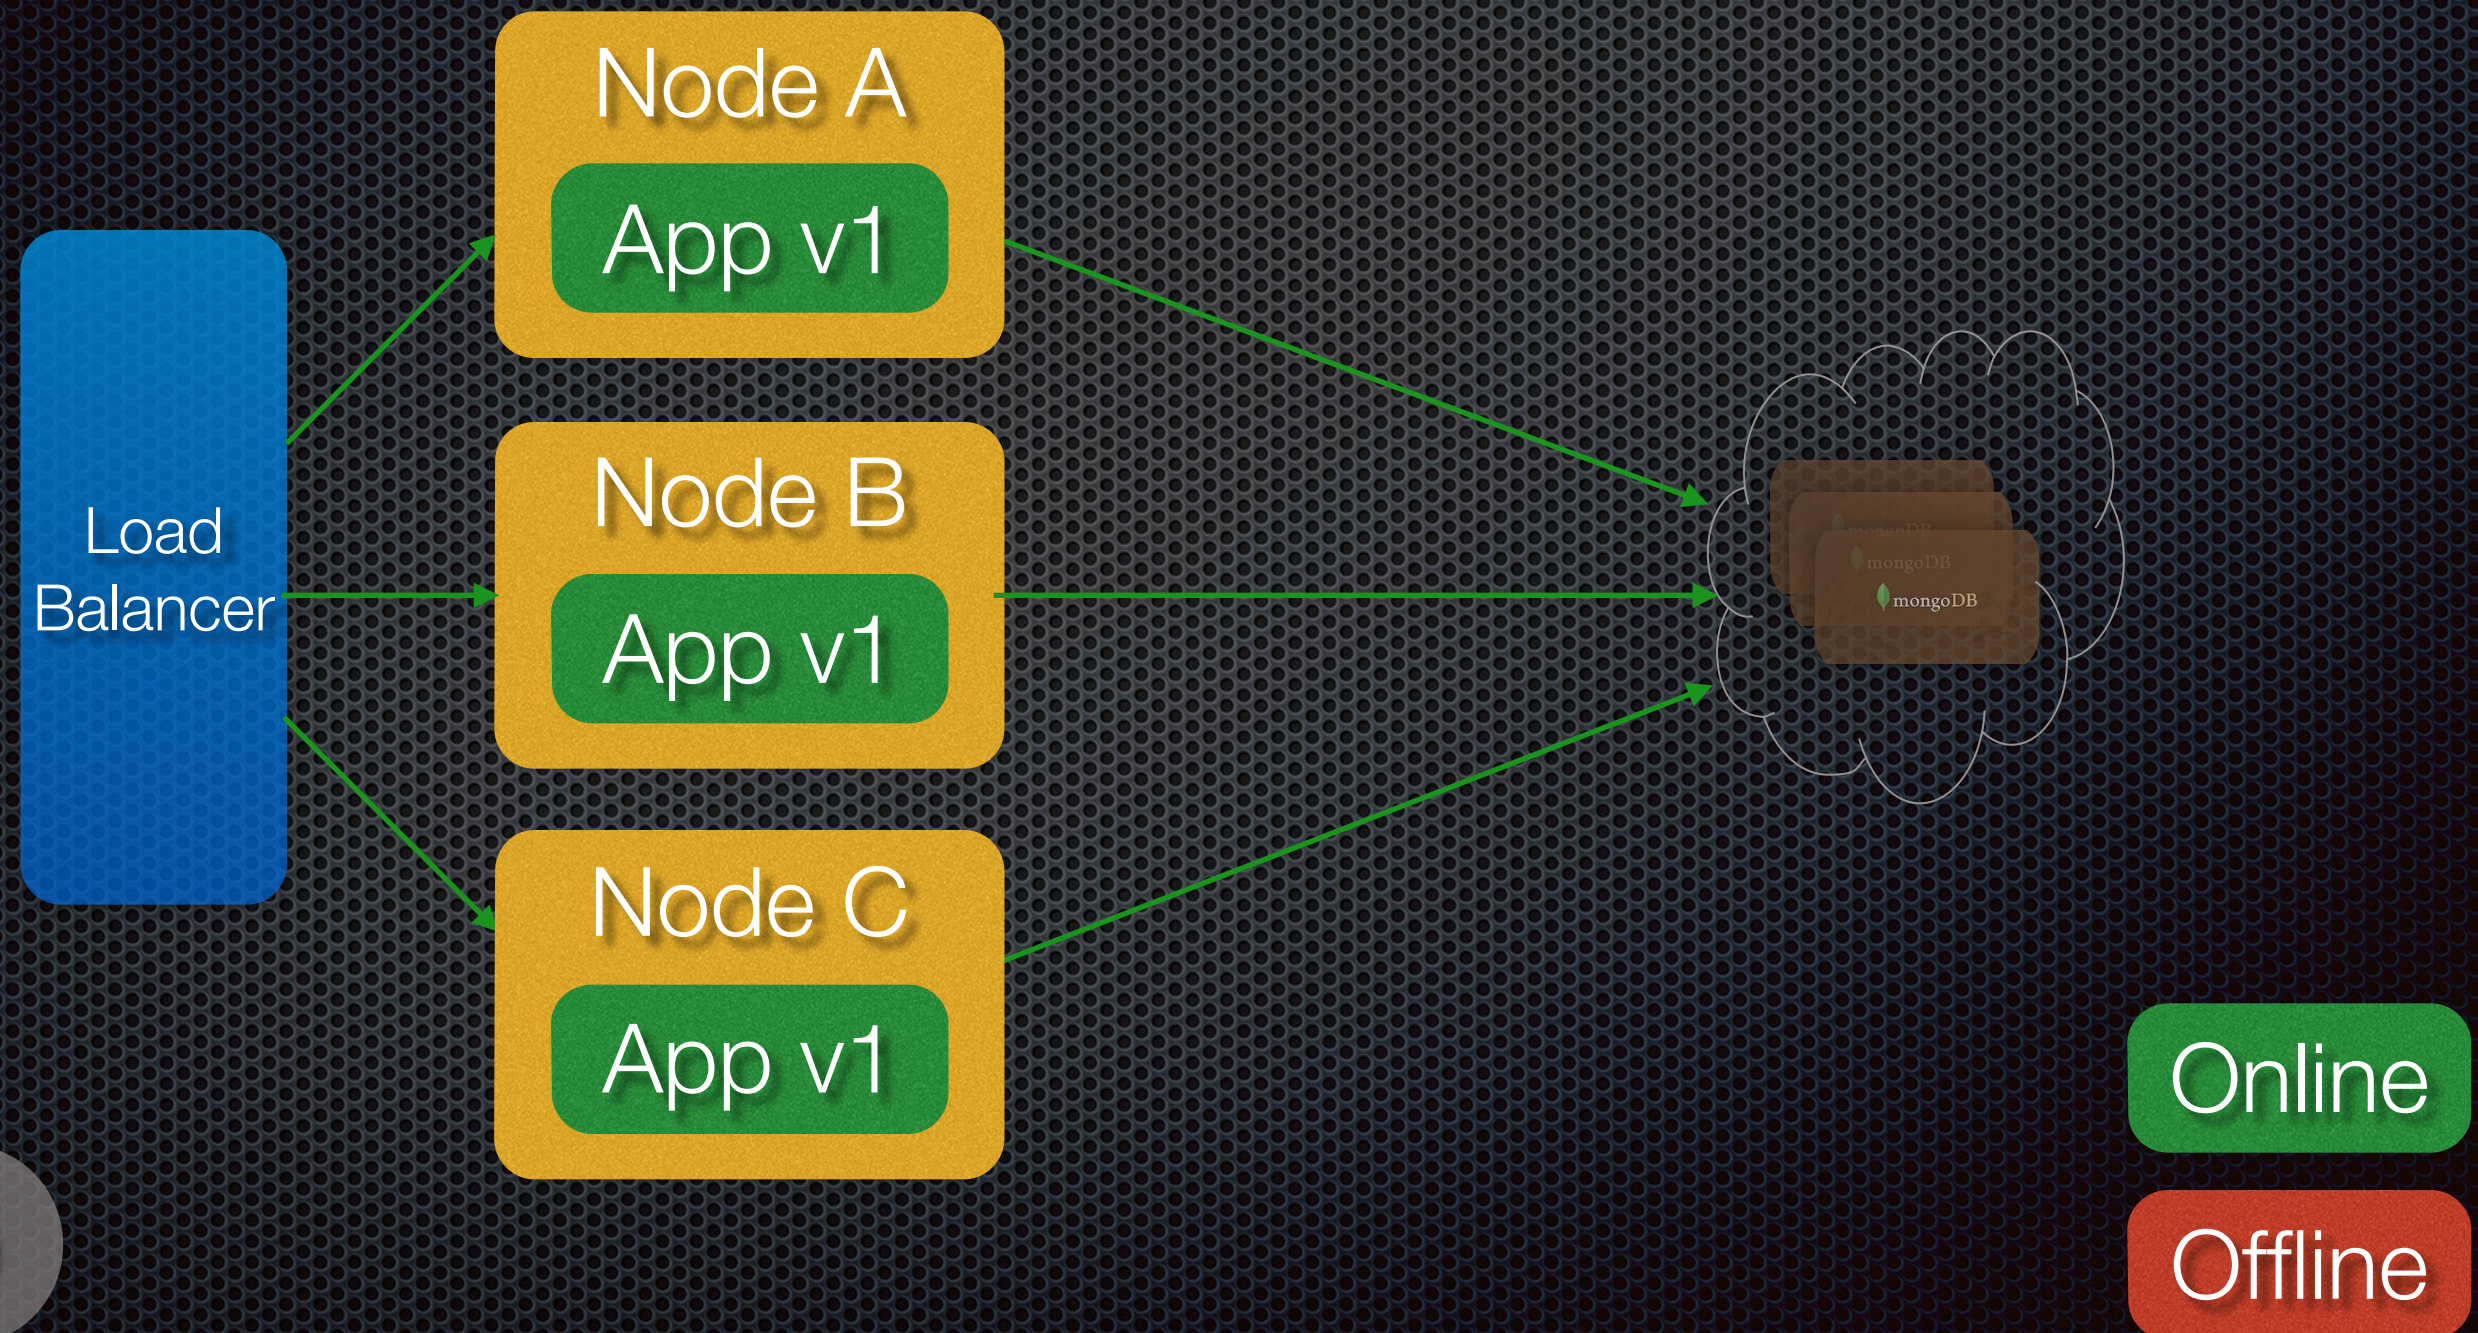
Deployment config - 2



Zero-Downtime Deployment Configuration - 3 AWS

Not Supported in this Release

Deployment config - 3



The Team

- ✦ Brian Blignaut
 - ✦ bblignaut@equalexperts.com
- ✦ Dhaval Dalal [@softwareartisan]
 - ✦ ddalal@equalexperts.com
- ✦ Vivek Dhapola
 - ✦ vdhapola@equalexperts.com
- ✦ Komal Jain
 - ✦ kjain@equalexperts.com

References

- ✦ Owen Rogers
 - ✦ <http://exortech.com/blog/2009/02/01/weekly-release-blog-11-zero-downtime-database-deployment/>

Thank-You